BEFORE YOU WALK IN · FRIDAY 23 OCTOBER

Corpus, Loop, Verifier

Thomas Joshi wrote one of these documents and is a co-author of the other. This is what both say, what to say back, and the nine facts to carry in your head.

DOCUMENT 1	The Neolab Wild West NEA, 17 June 2026
DOCUMENT 2	DSPy ICLR 2024 spotlight
THE POINT	You have two of his three layers. The one you have best is the one he calls most underestimated.
RULE	Use his framework. Don’t take his label.

Private. Prepared for Bhavit Saini. Every claim here is traceable to one of the two documents or to your own repository.

WHAT TO SAY, IN ORDER

01 The spine

Seven beats, about seventy-five seconds total. Learn the **order**, not the wording — the words will come out differently and that is fine. What must survive is the sequence.

#	BEAT	SAY THIS	IF HE INTERRUPTS
1	**The hook**	“Every knowledge worker has a second job they don’t get paid for: rebuilding what they were doing before the interruption.”	You already have him. Let him finish the thought.
2	**The product**	“Evo watches work happen locally, groups it into bodies of work, and puts the desk back when you say continue. Four seconds. No HTTP client in it.”	Skip the four seconds, keep the no-HTTP-client line.
3	**The concession**	“As a Mac utility that’s a five-to-forty-million-dollar business, and I’d be fine saying so.”	Never skip this. It licenses everything after it.
4	**The layer**	“The reason it’s bigger is what accumulates: not content, boundaries. Every agent today is handed its boundary — you point it at a folder. None derive it.”	This is the whole pitch. Come back to it.
5	**The bridge**	“Your neolab piece is three layers — corpus, loop, verifier. I have the corpus and the verifier. I skipped the loop, and I’d rather say that than pretend.”	He will lean forward here. Pause.
6	**The close**	“The verifier is the part I’d defend, because it’s the part nobody else in this category has.”	Then stop. Do not add a market size.
7	**The gift**	“And the switching-cost argument is already in your article: workflows embedded so deep that ripping them out costs more than keeping them.”	Use only if the conversation is warm.

The one thing to actually memorise

If everything else falls out of your head, three words survive: **corpus, loop, verifier.** They are his own framework, they describe your company honestly, and they give you a way to answer almost anything by pointing at one of the three.

> “By your framework I’ve built the corpus and the verifier and skipped the loop. The verifier is the one I’d defend.”

WHY IT WORKS

That single sentence does four jobs at once. It proves you read the article. It proves you can size your own work honestly. It names the weak layer before he does — which is the thing that makes an associate trust the strong one. And it gives him a sentence he can repeat to a partner, which is the actual currency of an associate meeting.

Page 3: the article. Page 4: the paper. Page 5: twelve questions. Page 6: numbers, never-says, fallback. Read 3 and 4 tonight; carry 5 and 6 into the room.

THE DOCUMENT HE WROTE

02 The Neolab Wild West

NEA's blog, 17 June 2026, by **Thomas Joshi, Madison Faulkner, Lila Tretikov and Andrew Schoen.** It is the fund's public thesis about where the next decade of AI returns comes from — a statement of what NEA wants to be true.

The argument, in one paragraph

Venture has left the frontier-lab era. **$1.68 trillion** of private AI-lab value exists and **93.6% of it sits in four companies.** Everything else combined is barely 6% — one basket, not a market. So the next returns come from elsewhere: **neolabs** — research-first companies that own a domain corpus, run a closed training-and-verification loop, and can be sold to a strategic.

Three trends, and how you stand on each

TREND	WHAT IT CLAIMS	YOUR POSITION
Commoditisation	Open weights close the gap; the marginal pre-training dollar buys less. "What separates top models is no longer what they know but how they act within specific environments."	Your intelligence is local deterministic code, not a model. You are downstream of this trend, not fighting it.
The budget shift	RL-first training moved the moat from compute to the loop: "compute is no longer the only moat; the entire loop is a moat."	Say **verifier**, not data. That is the word they chose.
The exit	Top-decile software trades at ~15–18x forward revenue against 3–6x for the rest. Strategics buy specific capability.	Your eventual buyer is an agent platform that needs a boundary it cannot assemble.

The three layers — this is the whole meeting

LAYER	WHAT THE ARTICLE SAYS	WHAT YOU ACTUALLY HAVE
Corpus	A domain-biased corpus pays back on every later cycle. Pre-train on one, then RL, beats RL alone — and the gap grows with difficulty.	**Real but small.** Six works, 164 labels, two days of one person's work. Nobody else can assemble it, but it is not a pre-training corpus.
Loop	"Whoever owns the environment owns the data flywheel." Every action inside it generates signal nobody outside can copy.	**Not built.** Deterministic core, no RL. A planned per-user fit on declarations — say that, and say it is not the same thing.
Verifier	"The most underestimated asset in this stack, because it looks like infrastructure when it is actually the source of truth."	**Strong, and genuinely rare.** Six assertions in the build. 164 labels of real work. Exit code 1 if two works ever fuse. Receipts on every restore.

WHY IT WORKS

You have been calling the gate an engineering chore. In NEA's vocabulary it is **the verifier layer** — the one they say is most underestimated.

The sentence to quote back

That middle clause **is** your switching-cost argument — they wrote it for you.

> "The durable wedges are domain-specific data that no one else can assemble, workflows embedded so deep in a customer's operations that ripping them out would cost more than keeping them, and distribution channels that the hyperscalers cannot easily replicate."

THE PAPER HE CO-WROTE

03 DSPy

What it is, plainly. Before DSPy, an LLM pipeline was a pile of prompt strings someone tuned by hand until it worked. DSPy says: stop writing prompts, write a program.
Declare each step — "takes a question, returns an answer" — wire them together, and let a compiler search for the instructions and worked examples that maximise **a metric you defined.** The prompt stops being something you hand-craft and becomes a parameter the compiler finds.

	WHAT TO KNOW
The paper	DSPy: Compiling Declarative Language Model Calls into State-of-the-Art Pipelines. **ICLR 2024, spotlight.** First author Omar Khattab; senior authors include Matei Zaharia and Christopher Potts. **Thirteen authors; Thomas T. Joshi is the ninth.**
The claim	Within minutes of compiling, it beats out-of-the-box few-shot prompting and expert-written demonstrations on GPT-3.5 and Llama2-13b. A compiled 770M T5 becomes competitive with large proprietary models driven by hand-written prompt chains.
NEA's line	They call it "the most popular open source software to come out of Stanford AI," used by Meta and Microsoft. His bio says **co-author.** Not "the creator."

The bridge — the best thing you can say to this man

DSPy's real idea is not the compiler. It is that **you do not hand-tune the system; you declare what must be true and let the build be measured against it.** Your gate is the same move, applied to something deterministic.

SAY THIS

"The thing I took from DSPy isn't the library — there's no language model in my restore path. It's the discipline. You don't hand-tune the grouping; you declare the properties that must hold and let the build fail when they don't. Six assertions: corpus is real, labels are real, something is reconstructed, no two works fuse, no leisure resource ever anchors a thread, same input gives the same output twice. Same idea, deterministic system."

WHY IT WORKS

True, specific, checkable in his head, and it connects his own published work to yours without flattering him. Associates get flattered constantly. Almost nobody says "here is how your paper changed how I build the test that fails me."

Three traps

TRAP

Don't say "you wrote DSPy," and don't claim you use it. He is one of thirteen authors; NEA's own bio says **co-author**, which is correct and impressive — "creator" is wrong and he will correct you. And there is no LM in your loop, so if you say "we're doing DSPy for work grouping," his next question is "which teleprompter?" and you have nothing.

TRAP

Don't say "we do RL." You don't. Say **deterministic core plus a planned per-user fit.** Different words, and he will know within one question.

ONE LINE EACH

04 Twelve questions

Left column is what he says. Right column is the answer — short on purpose. If you find yourself extending an answer past two sentences, stop and ask him a question instead.

IF HE ASKS	ANSWER
“Did you read the neolab piece?”	“Twice. The three-layer bit is the reason I built the gate the way I did.”
“What did you take from it?”	“Corpus, loop, verifier. I have two of the three and I’d rather tell you which one is missing.”
“Which one is missing?”	“The loop. My core is deterministic. The learning I’d add is a small per-user fit on declarations — that’s not an RL loop and I’m not going to call it one.”
“Then you’re not a neolab.”	“No. I’m a product company in a neolab-shaped thesis, at pre-seed. Calling myself a neolab at $1.5M would be silly.”
“Where’s your verifier?”	“A corpus of 164 labels across six real works from one person’s two days, and six assertions that fail the build. If two works ever fuse, exit code 1.”
“Isn’t that just a test suite?”	“A test suite checks that the code runs. This checks that the boundary is still correct — it’s six properties, not six behaviours.”
“How big is the corpus really?”	“Ten thousand events, two days, one person, six works. It’s small. I know exactly how small.”
“Who owns the corpus?”	“The user. It never leaves the machine — there’s no HTTP client in the build. That’s the whole reason the layer has to be neutral.”
“What’s DSPy, in one line?”	“Declare the steps, define the metric, let a compiler find the prompt. The metric is the program.”
“Do you use it?”	“No — no language model in the loop. I took the metric discipline, not the library.”
“What would you compile if you could?”	“The boundary, not the safety. Learned layer on top, deterministic core underneath, and the core is the part that can’t be argued with.”
“Why won’t OpenAI build this?”	“They did — Computer History, August. It’s ungrouped, because their unit is the conversation. What they can’t add is neutrality: a layer that lets Claude be safe on your desktop is against their interest.”

The answer that wins the room

If he gives you one sentence to defend, make it this one — it is the only claim you can make that nobody else in the category can.

> “Every layer in that piece depends on someone owning the environment. My user owns it, not me. That’s why the corpus can be neutral, and neutrality is the only thing that makes a jurisdiction layer worth anything.”

NUMBERS, NEVER-SAYS, FALLBACK

05 The card

Numbers worth carrying

NUMBER	WHAT IT IS
$1.68T	Aggregate private valuation of AI research labs, April 2026 (PitchBook, via NEA).
93.6%	Share sitting in four frontier generalists. Everything else combined: ~6%.
15–18x	Forward revenue multiple for top-decile software. The rest of the index: 3–6x.
$50M–$500M	What neolabs raise before revenue. You are raising $1.5M — hence the label trap.
86 labs	Companies in NEA's map. Not one is a jurisdiction layer.
13 / ICLR	Authors on the DSPy paper; Thomas is ninth. Accepted as a **spotlight**.
164 labels	In your verifier corpus. Six works. Two days. Exit code 1 on fusion.

Never say · Always say

NEVER	ALWAYS
"We're a neolab."	"A neolab-shaped thesis at pre-seed."
"You created DSPy."	"Co-author of DSPy."
"We do RL."	"Deterministic core, planned per-user fit."
"We use DSPy."	"I took the metric discipline, not the library."
"The corpus proves it."	"The verifier is the part I'd defend."
"No one is doing this."	"Everyone is doing the feature. Nobody is doing the layer."
"Our AI pipeline…"	"Deterministic code, with a small model for transcription."

The fallback — for any question you cannot answer

> **IF YOU ARE STUCK**
> "I don't know. That's a real gap, not a hedge. I know the shape of the answer and how I'd test it. Can I send it by Wednesday — and would you tell me what would actually change your mind?"

> **WHY IT WORKS**
> Four beats: admit, nearest true thing, a date, turn it back. The last clause is what makes him remember you — it treats him like a researcher rather than a gatekeeper. He is a researcher.

One last thing. You do not need to remember this sheet. Three words — **corpus, loop, verifier** — and one sentence: "I have two of the three, and the verifier is the one I'd defend." Everything else is here so you recognise the question when it arrives.

Sources: nea.com/blog/the-ai-neolab-wild-west (17 Jun 2026); openreview.net/forum?id=sY5N0zY5Od (ICLR 2024 spotlight); nea.com/team/thomas-joshi. Numbers are quoted as the article states them, with its own PitchBook and Epoch AI citations.