PRIVATE · PREPARED FOR BHAVIT

The NEA Meeting

Everything you need for the thirty minutes that decide whether this is a company.

MEETING	Thomas Joshi, Associate — NEA
WHERE	Friday 23 October 2026 · 11:00am Pacific Zoom · thirty minutes · 11:30pm your time
CAME VIA	Lila Tretikov, Partner, Head of AI Strategy
THE ASK	$1.5M on a $12M post-money SAFE

Prepared for Bhavit Saini. Private and confidential. Every claim in this book carries a confidence grade and a source — see Section 23.

Contents

PART I — ORIENTATION

00 How to use this book 3
01 The thirty minutes, minute by minute 4
02 What changed in the code 5

PART II — THE COUNTERPARTY

03 Thomas Joshi, in detail 7
04 NEA: the machine behind him 9
05 The memo he has to write about you 11

PART III — THE HONEST COMPANY

06 The claim ledger 13
07 The traction conversation 16
08 The vibe-coding answer 18
09 What the code actually does 20
10 Incorporation and visas 24

PART IV — THE MEETING

11 The first four minutes 26
12 The demo 28
13 The forty questions 30
14 The ten that decide it 37
15 When you don't know 38

PART V — THE CLOSE

16 Your questions for him 41
17 The ask, the cap, the negotiation 43
18 After the call 46
19 If it's a no, and if it's a yes 48

PART VI — REFERENCE

20 The one-page card 49
21 Numbers you must know cold 50
22 Do not say 51
23 Sources and confidence 52

Sections 00-05 are the ones to read tonight. Section 05 explains the whole meeting in three pages. Section 20 is the card you keep next to you on the call.

START HERE

00 How to use this book

This book exists for one meeting: **Friday 23 October 2026, 11:00am Pacific, thirty minutes, Zoom, with Thomas Joshi, Associate at NEA.** Nothing else. It is not a general fundraising guide and it does not cover any earlier project of yours.

You are a first-time founder and you have two days. So the book is built backwards from a single question: what does Thomas need to believe at 11:30am on Friday that he does not believe at 11:00am? Everything here serves that.

Read it in this order

1. **Section 01 and 20.** The run of show and the one-page card. Read these first, tonight. They tell you what the thirty minutes actually look like.
2. **Sections 03, 04 and 05.** Who Thomas is, what NEA is, and — the most important section in the book — the memo he has to write about you after the call. If you understand the memo, every answer in the meeting becomes obvious.
3. **Sections 06 to 10.** The honest company: what your deck overstates, what the code actually does, how to say the vibe-coding thing, and the visa answer.
4. **Sections 11 to 15.** The words. Opening, demo, forty questions, the ten that decide it, and what to say when you don't know.
5. **Sections 16 to 19.** The close, the ask, and what happens afterwards.

Four rules for the thirty minutes

Rule	What it means in practice
Never invent a number	Not once. One invented figure contaminates every true one, and an associate's job is to check. If you don't know, say so — Section 15 gives you the exact words.
Short answers	Twenty to forty seconds. He will ask a follow-up if he wants more. Long answers read as anxiety and eat the clock.
Correct your own deck before he does	Every overstatement in your deck is a trap you set for yourself. Section 06 lists them and gives you the sentence that defuses each.
Ask for the next step out loud	He will not offer one unprompted. The last ninety seconds are yours and most founders waste them.

Two layers, everywhere

Every answer in this book comes in two layers: **one sentence you can say out loud** , then the deeper layer if he leans in. You do not need to memorise the deeper layer. You need to know it exists, so that a follow-up doesn't knock you over.

> **WHY THIS WORKS**
> You are not being tested on being impressive. You are being tested on whether you are legible, honest, and specific. Thomas has met hundreds of founders. He cannot tell in thirty minutes whether Evo will win. He can tell, very quickly, whether you are the kind of person who knows what he knows and what he doesn't. That is the whole game.

RUN OF SHOW

01 The thirty minutes, minute by minute

Thirty minutes is short. It is roughly **four thousand words** of conversation in total, both sides. If you spend the first six minutes on a monologue, you have spent a fifth of the meeting.

Assume this shape. It will not go exactly like this — he may open with questions — but knowing the shape stops you from panicking when it drifts.

Clock	What is happening	Your job
0:00–0:02	He says hello, thanks you, maybe apologises for rescheduling. Small talk.	Be warm, be brief. Do not fill silence.
0:02–0:05	"Tell me about Evo." This is the only open invitation you will get.	Deliver the four-minute opening in Section 11, verbatim or close to it.
0:05–0:12	Either a demo, or he starts drilling.	Offer the demo explicitly: "I can show you in three minutes, or keep talking — your call." Then follow his answer.
0:12–0:25	Questions. This is where the meeting is won or lost.	Sections 13 and 14. Short answers. Concede fast where you're weak.
0:25–0:29	Your turn.	Ask three questions from Section 16. The first one is always what happens next and on what timeline.
0:29–0:30	Close.	One sentence on what you'll send. Then stop talking.

The only thing you must not do

Do not let the meeting end without knowing **what the next step is and what would make it a yes.** If you leave and cannot answer those two questions, the meeting failed regardless of how well the conversation went. Section 16 gives you the exact phrasing so it doesn't sound like begging.

> **CAREFUL**
> The meeting is at 11:00am Pacific on Friday 23 October. That is 11:30pm IST. If the call is at 11:00am PT and the window is an hour, you will be on the call from 11:30pm to midnight your time. Eat before. Turn every notification off. Do not take it from a bed or a moving car. Section 20 has the physical checklist — read it the night before, not the morning of.

Before Friday, do these five things

1. **Confirm the exact start time and ask for the Zoom link** if you don't already have it. Two lines, no paragraph. He has asked for times twice already; he responds to short emails, so keep it short.
2. **Run** `cargo test --workspace` once. **You will be asked, directly or obliquely, how the codebase is. You want a number from this week, not from September.**
3. **Record a 90-second screen demo** with narration, and put it in a link you can paste in one line. It is your fallback and your follow-up asset.
4. **Write the three sentences you will say about traction** (Section 07). Say them out loud until they don't sound defensive.
5. **Read Section 05 twice.** The memo is the whole book compressed.

SINCE YOUR LAST PREP

02 What changed in the code — and why it changes the conversation

Your repo moved a long way between late September and mid-October. That matters for this meeting, because **the strongest version of your story is no longer "I built an app that restores windows."** It is "I built an engine that reports what it actually did, and refuses when it cannot be sure." That is a very different, and much more investable, sentence.

What is actually new

Surface	What it does now	Why an investor should care
Pods	A body of work became a first-class object you switch between: a card with a name, a colour, a mini-stage, a one-line reason for where it resumes, and hotkeys. Entering one arranges its windows into a layout the work remembers across days.	This is the first visible product surface that is not a list. It shows the model, not the machinery.
Restore receipts	Every restoration now reports per-target, in three buckets — opened, raised (brought an existing window forward), or unopened-with-reason — plus a separate verification result: Verified, Unverified, Failed.	It means the product can tell you what it didn't do. Almost no consumer product does this. Thomas's whole thesis is about evaluation; this is evaluation shipped as UI.
The ground-truth gate went hard	A committed corpus from one real recorded day — 10,392 events, 164 hand-labelled resources across six bodies of work — replays through the engine inside cargo test, with hard assertions: no thread may fuse two labelled works, no leisure resource may anchor a work, and replay must be deterministic.	It converts "we have a model" into "we have a test that fails when the model is wrong." This is the single most impressive thing in the repo to a technical investor.
Gestures are declarations	Dragging a window onto a pod card, merging two cards, ejecting a window — each emits the engine's own Declaration. The physical correction is the merge path. There is no code path from inference to a merge.	It answers "what happens when it's wrong?" with architecture instead of promises.
Voice that reads the engine	ASK and BRIEF verbs narrate the actual engine state — "You made Evo desktop. You left off at Evo desktop." — computed at request time, so the spoken answer is as fresh as the log.	The voice is not a chatbot bolted on. It is a read-out of the same evidence the rest of the product uses.
The restoration audit	An engineering pass that found and fixed real defects: a wrong-window risk on duplicate titles, a voice path reading a different authority than the read path, an embedding loader that could reach the network.	It exists as a written document in the repo. If he asks "how do you know it works?", you can point at a document that lists what does not.

Three numbers to carry in your head

16 crates — Rust workspace, roughly 138,500 lines. Native daemon, native desktop shell, no Electron.

~1,270 tests — 1,272 passed, 0 failed, 4 ignored, across 59 suite results at the last recorded full offline run (5 October 2026). The September figure was 1,087 across 57 suites. Re-run it before Friday so you can say this week.

1 real corpus	10,392 events, 164 hand labels, six works, two days of one person's actual machine. Small. Say it's small.

CAREFUL

Do not say "1,087 tests" on Friday if the repo says 1,272. Do not say "CI runs on every commit" — there is no .github directory in the repo. What is true is that the gate is a hard test inside the workspace suite, and that the same check exists as a CLI that exits non-zero. Say that. Section 09 gives you the exact phrasing.

The one-paragraph version of the new story

SAY THIS

"The thing I'd want you to see is not the window count. It's that the engine has to be honest. Every restore now reports what it opened, what it brought forward, and what it couldn't open — three buckets, no silent skips. And there's a corpus of a real working day in the test suite that fails the build if two separate bodies of work ever merge. I'd rather ship a product that tells you what it didn't do than one that looks seamless and lies."

WHY THIS WORKS

This paragraph does something no amount of enthusiasm can: it makes you sound like an engineer to an engineer, without claiming to be one. It is also completely true, which means it survives any follow-up.

03 Thomas Joshi, in detail

He is an Associate. That word matters more than founders realise, and it cuts both ways. He cannot write you a cheque on his own. But he is the person who decides **whether the partners ever hear your name** — and in a firm that runs on partner time, that is real power.

The facts on the record

Role	Associate, Technology Investing Team. Joined 2025. Covers seed / early stage and growth.
Background	AI researcher and AI engineer before investing. Built infrastructure supporting language and classification models for conversational AI in healthcare.
Education	B.Sc. in Artificial Intelligence, Columbia University.
Known for	Co-author of Stanford DSPy — the prompt-programming framework; one of the most widely used pieces of research software to come out of Stanford AI.
Published	Co-wrote "The AI Neolab Wild West" (17 June 2026) with Lila Tretikov, Madison Faulkner and Andrew Schoen. Co-wrote the CuspAI follow-on post (23 July 2026).
Deals on his page	Anyscale (first invested 2019), Metabase (2019), Factory (2025). Led Namespace Labs' Series A; NEA participated in their $42M Series B announced late September 2026, led by Scale Venture Partners.
Also public	Organised a Neolab event at NEA with Snowflake, Factory and Blackstone, where he gave a deep dive on recursive self-improvement. Hosts a podcast. Active on X.

The five things he has actually said

These are worth learning close to verbatim, because they are his words and using them back to him is not flattery — it is evidence that you did the reading. Do it once, lightly, and never twice.

> The utility problem in AI. We have a working recipe for capability, but the second half is about defining the right problems and owning evaluation so products create real economic value.
>
> — Thomas Joshi, NEA profile

> We need to treat evaluation, task design, and human-in-the-loop feedback as its own new layer.
>
> — Thomas Joshi, NEA profile

> Living benchmarks, longitudinal memory and agency, and outcome-grounded metrics.
>
> — Thomas Joshi, NEA profile — his description of what the field needs

> You can tell when a product has been infused with love.
>
> — Thomas Joshi, NEA profile

> We led their Series A after identifying an emerging bottleneck in the software development lifecycle, and knew that [the founder] was the perfect person to address it.
>
> — Thomas Joshi, on Namespace Labs, X, October 2026

Why those five quotes are a gift

Look at what they add up to. He is not a momentum investor. He is not looking for the biggest TAM slide. He is looking for **a bottleneck in a real workflow, a person unusually positioned to remove it, and evidence that the removal can be measured.**

That is almost a description of Evo. Your bottleneck is context reconstruction. You are unusually

positioned because you lived the problem and built the engine. And the measurement exists — it's the corpus and the recovery buckets.

> **SAY THIS**
> "The way I'd frame what we're doing, in your language: the bottleneck in knowledge work isn't intelligence any more, it's context. The models are good enough; they just don't know where you were. That's an emerging bottleneck in a workflow, and it's measurable — either the workspace came back or it didn't."

> **WHY THIS WORKS**
> "Emerging bottleneck" and "measurable" are his own phrases from his own post about Namespace. You are not quoting him at him; you are showing that you understand what he looks for. Say it once, in the opening, then never again.

The uncomfortable read

Now the part nobody will tell you. Look at the deals above: Anyscale, Factory, Metabase, Namespace, CuspAI. That is **infrastructure for AI applications, developer tooling, and frontier research.** Enterprise or deep-tech, in almost every case.

Evo is a **consumer Mac app** with a prosumer price point. On paper, that is not his pattern. There are three honest readings of why he took this call anyway, and you should hold all three at once:

1. **Lila forwarded it.** She is Head of AI Strategy and she wrote the Neolab post. A partner forwarding your email is the single strongest asset you have in this meeting. It is also not the same as conviction.
2. **NEA is actively building a memory thesis.** NEA's own fellowship write-ups describe fellows working on exactly that: "the AI memory thesis and building out a people intelligence layer." Evo is a memory company. You are on-thesis even if you are not on-pattern.
3. **He is curious and it cost him nothing.** Associates take thirty-minute calls. It is cheap optionality. Do not read enthusiasm into politeness, and do not read politeness as a no either.

> **CAREFUL**
> The most likely failure mode of this meeting is not that you say something wrong. It is that you spend thirty minutes being impressive about a prosumer Mac utility, and he writes a memo that says "nice product, too small for us." Section 05 is the antidote, and it is the most important section in this book.

04 NEA: the machine behind him

You need to understand one thing about NEA before you decide how to pitch: **the arithmetic of a very large fund.** Not because it is fatal, but because it explains every strange thing that might happen in the meeting.

The scale

Founded	1977. One of the oldest venture firms in the world.
Assets	More than $38 billion under management as of 30 June 2026. $35 billion at 31 December 2025. The figure moves; quote it as "roughly $38 billion" if you must quote it at all.
Current fund	NEA 19, roughly $5.5 billion target, first close around December 2025 – January 2026.
Track record	More than 285 portfolio company public listings and 515+ mergers and acquisitions, on the firm's own count.
Stage	Seed through IPO. Historically very IPO-heavy for a venture firm.
Typical seed round at NEA	Tracxn counts 241 seed investments at an average round size of about $4.4M. That is the round, not the cheque.

The arithmetic that nobody says out loud

Run this once, privately, and then forget the number:

- $1.5M into NEA 19 is **0.027% of the fund.**
- If Evo returned **100×** , NEA's $1.5M becomes $150M — about **2.7% of the fund.**
- To return the whole fund, that single $1.5M cheque would have to become **$5.5 billion.**

This is why large funds often pass on genuinely good small companies, and it is why the phrase "too early for us" usually means "too small for the fund." It is not an insult and it is not about you.

But — and this is the part founders miss — **it does not kill you.** Firms this size routinely write small cheques for three reasons: to keep a relationship warm, to buy an option on a category, and because a partner championed it. Your job is to give Thomas a reason that is one of those three, and to be the obvious choice if the category turns out to be real.

> **SAY THIS**
> "I know a $1.5M round is small for NEA 19. I'm not asking you to make it the fund's return. I'm asking whether the category is one you want a position in before the seed, and whether we're the right team to hold it."

> **CAREFUL**
> Do not open with that. It is a closing argument, not an opening one. It works at 0:25, when he has already decided you are interesting. Said at 0:05, it is a reason to say no.

The process, as best it can be reconstructed

Step	What happens	What it means for you
1. Associate call	This meeting. Screening. No decision.	Goal: he wants to keep talking. Nothing more.
2. Deeper diligence	Two to six further calls, a technical session, references, sometimes a customer call.	Goal: survive the boring questions.
3. Partner meeting	Monday partnership meetings are part of NEA's long-standing rhythm. Partners vote — historically yes, yes-with-reservations, or no — with a second vote later.	Goal: arm Thomas with the memo he can defend in that room.
4. Term sheet and close	Legal, then wire.	Goal: don't celebrate until the wire lands.

Realistically, from a first associate call to money is **four to eight weeks** if it happens at all. Plan for that. Do not stop fundraising, and do not spend the money before it lands.

What "on-thesis" means at NEA right now

Three current threads you should be able to speak to, because they are the water he swims in:

- **Vertical AI and workflow data.** NEA's stated AI framework leans toward industry-specific agents that process unstructured information and integrate with existing business software, where newly captured workflow data becomes the advantage. Your per-user work graph is workflow data captured at the source.
- **The Neolab thesis.** Proprietary domain corpus + closed loop + domain-specific verifier. You are **not** a neolab — don't pretend to be — but the three-layer logic maps onto Evo one level up: the corpus is the per-user evidence graph, the loop is restore-and-correct, the verifier is whether the person kept working in the restored workspace.
- **Evaluation as a product problem.** This is his most consistent theme. Your recovery buckets and the ground-truth gate are the direct answer, and almost no pre-seed company has them.

> **NOTE**
> A small tactical point that punches above its weight: NEA's own fellowship cohort has been working on an AI memory thesis and a people intelligence layer. That means there is probably an internal document somewhere that already argues your category matters. You cannot cite it, but you can speak to the shape of it. If he says "memory is a crowded space," you can say: "It is — but almost all of it is memory of conversations. Almost none of it is memory of work."

05 The memo he has to write about you

Here is the single most useful thing in this book. After the call, Thomas will write a short internal memo. It will be two or three paragraphs. If it is good, partners read it. If it is bad, nothing else you did matters.

You cannot see it. But you can predict its shape, because memo-shaped writing is formulaic. It will answer five questions in roughly this order.

The five lines, and what decides each

Line in the memo	What fills it	How you win it
What is it?	One sentence. If he can't write one sentence, you've lost.	Give him the sentence yourself, early, and repeat it at the end. "A memory layer for work that restores your whole workspace when you say continue."
Why now?	What changed recently that makes this possible or necessary.	Three things, concretely: on-device models got small enough; capture APIs matured; and an OpenAI product made the category legible. Section 13 gives you the exact phrasing, including the honest caveat.
Why this team?	Is there a reason these people win, or are they interchangeable?	Not "I'm technical." The honest version: you lived the problem for two years, you designed a behaviour that is genuinely hard to copy, and you know exactly which parts you did not build. Section 08.
Is it working?	Evidence, graded. He will discount enthusiasm to zero.	The corpus. The recovery buckets. The gate. And the honest traction number, which is smaller than your deck says. Section 07.
Is it big enough?	This is the one you are most likely to lose.	Not TAM. The path from a $108 utility to the context layer that agents query before they act. Section 13, question 31.

The best possible memo

Imagine the paragraph Thomas wishes he could write. It looks like this:

> Pre-seed, $1.5M at $12M post. Founder built a local-first work memory for Mac — it watches app/file/tab activity on-device and reconstructs the whole workspace on command. Unusually for pre-seed, he has an evaluation harness: a labelled corpus from a real working day, hard regression gates against work fusion, and per-target restore receipts. Traction is early and verbal, but the product is real and he uses it daily. Category just got validated by OpenAI's Computer History. Risk: prosumer price point and single founder. Angle: this is the context/jurisdiction layer for agents, which is our Namespace/Factory thesis one layer down. Worth a partner meeting.
>
> — What you are aiming for — not a quote from anyone

Read that again. Notice what is **not** in it: no TAM number, no "10x better," no adjective about how big this could be. It is specific, it concedes the risk, and it names a thesis the firm already holds. That last move is the one that gets you a second meeting.

The worst possible memo

> Consumer Mac utility that records what you do and reopens your windows. Solo founder, pre-revenue, no written traction. Competing with OpenAI and Microsoft. $12M post feels rich for a feature. Pass.
>
> — What you are trying to avoid

Everything in that paragraph is either true or arguable. That is what makes it dangerous. Your job on Friday is to make the first paragraph easier to write than the second — and the way you do that is by **naming the risk yourself, first, and then showing the path through it.**

WHY THIS WORKS

This is counter-intuitive and it works almost every time. A founder who says "the honest risk here is that this looks like a feature" is treated as an operator. A founder who waits to be accused of it is treated as a salesperson. Same fact, opposite outcome.

The one sentence to leave in the room

If he remembers one line on Monday morning, you want it to be this:

SAY THIS

"Every agent is getting memory. Almost none of them are getting jurisdiction — knowing which of your things belong to which job. That's what we produce, and we produce it as a side effect of being useful every morning."

It is short, it is his vocabulary (memory, agents), it reframes you away from "window restorer," and it is true. Say it once, near the end, and then stop.

EVERY NUMBER IN YOUR DECK, GRADED

06 The claim ledger

Your deck has eleven slides and roughly twenty-five factual claims. Most are defensible. Five are not, and one of those five will end the meeting if he catches it before you do. This section grades every one of them, so that you are never surprised.

The grades: **A** — you can defend it with a source or a measurement. **B** — directionally true but softer than it sounds. **C** — you should reframe it before he asks. **D** — do not repeat it on the call.

The deck, slide by slide

Slide	Claim	Grade	What to say instead
Problem	25 min 26 s average time before people return to interrupted work (Mark et al., CHI 2005)	A, with a caveat	Keep it, but know what it measures: the average time before a worker returns to the original task, during which they do other work. It is not literally time spent reconstructing. If he pushes: "It's the return-to-task figure, not a stopwatch on re-orientation."
Problem	~1,200 app toggles per workday (HBR 2022)	A	Fine. It's 137 users across 20 teams in 3 Fortune 500s, and the paper puts reorientation at just under 4 hours a week. Don't over-quote it.
Problem	Interrupted every 2 minutes (Microsoft WTI 2025)	B	It's telemetry from Microsoft 365, and the 2-minute figure is for the top 20% of users by notification volume. Say "up to" or "for the heaviest-notified users."
Product	MVP running today on macOS	A	True. Add the honest qualifier: "It's a working build I use; it's not a signed, notarised public release."
How it works	Everything on your device. Everything encrypted.	C	"Everything runs on the device. The local store is protected by FileVault and the app sandbox today; per-record encryption with a Keychain key is required before sync ships." Don't claim encryption you haven't built.
How it works	learn their state / "classification engine is our core IP"	B	The grouping is real and evidence-based. The learning is not — there's no fitted model yet. Say "the engine accumulates evidence and obeys corrections; fitting per-user weights is the next layer." See Section 09.
Traction	300+ organic waitlist sign-ups	D	There is no list. Say: "I've spoken to 300-plus people at hackathons, college events and through work — PMs, founder's-office people, SDEs, strategy analysts. It's verbal, not a list. I can name where and when."
Traction	26 MVP testers running it on their work machines	D	Nobody has run the current build on their own machine. Say: "26 people used an earlier build; 19 said they'd use it daily once it's finished. Nobody's on the current build yet."
Traction	19 of 26 still active at day 30 / D7 22/26, D30 19/26	D	This is the danger. It was a previous build, informal, measured by whether people came back — not instrumented. Say: "That was an earlier build, measured informally. I'm not going to present it as retention data."
Traction	50+ user interviews	B	True in substance if you can name who and where. Be ready with two specifics: a role, a place, and the sentence they said.
Traction	Three user testimonial quotes	D	They're composites written to convey what people said about the concept. If asked: "Those are paraphrases of what people told me, not verbatim quotes from daily users." Never present them as logged feedback.

Slide	Claim	Grade	What to say instead
Why now	Apple ships a 3B model to every Apple-silicon Mac	C	Approximate. Safer: "Apple ships on-device foundation models through a public framework, which wasn't true two years ago." Don't quote the parameter count.
Why now	40% of enterprise apps ship task-specific agents by end-2026 (Gartner)	B	It's a Gartner projection, not a measurement. Say "Gartner projects" and move on.
Why now	Microsoft Recall's privacy revolt, Rewind shut down	C	Rewind didn't shut down — Meta acquired Limitless in December 2025 and the Mac app was discontinued. Recall didn't fail; it's shipping, Windows-only. Say: "Recall launched into a backlash and stayed Windows-only; Rewind's Mac app was discontinued after the Meta acquisition. There's no trusted mass-market answer on Mac."
Market	6.1M / 70.7M / ~1B	A	BLS 2024 for the first two; the ~1B is a Gartner-scale estimate. Label the last one as an estimate.
Competition	ChatGPT Computer History — cloud, EU/UK-blocked	C	It launched EU/UK-blocked in August 2026 and was extended to the EEA, Switzerland and UK for Enterprise in September. Correct it before he does: "It's now in Europe. What hasn't changed is that the raw event stream is processed on OpenAI's servers."
Competition	Everyone remembers. Nobody restores.	B	Strong line, mostly true, but Microsoft Recall does have a "jump back in" feature for Office. Keep it as positioning, not as fact.
Business	Pro $108/yr	B	A hypothesis. No one has paid. Say "that's the hypothesis" if asked.
Business	Free to trust, $9/month to live in	B	Good line. Nothing behind it yet.
GTM	5,000 activated users by week 6	C	This is a target, not a plan with a budget attached. If asked how: name the channels and the assumption. Do not defend the number as a forecast.
Ask	$1.5M at $12M post-money SAFE, 18 months runway	A	Arithmetic. 1.5 ÷ 18 = $83K/month. Know that he will do this division in his head — see Section 17.
Ask	Seed from month 12	C	That means you start raising a seed with roughly six months of cash left, on the strength of 100K users and ~$1M ARR. It is a plan, not a forecast. Don't oversell it.
Team	Built Evo's Rust + SQLite engine	D as written	You designed it and directed it; you did not hand-write 138,000 lines. Section 08 gives you the full answer. Do not let the deck's phrasing stand unchallenged — correct it yourself.
Team	Buildspace S5 / B.Sc. CS, Delhi / Manavv's background	A	Your own history. Fine.

The three sentences that fix almost all of it

You do not need to memorise the table. You need three sentences, deployed at the right moment. The trick is to say them **before** he asks, because a correction you volunteer reads as rigour and a correction he extracts reads as a caught exaggeration. Same words, opposite meaning.

> **SAY THIS**
> "Before we go further — two things in that deck are softer than they look, and I'd rather say them now."

Then: the waitlist is verbal, the retention number is from an earlier build, the quotes are paraphrases, and the engine doesn't fit its own weights yet. Four sentences, twenty seconds, and the temperature of the meeting drops in your favour.

WHY THIS WORKS

This is the highest-leverage move available to you on Friday, and almost nobody does it. An associate's default assumption about a pre-seed founder is that the numbers are inflated. When you pre-emptively correct your own deck, you invert that assumption: now everything you don't correct is presumed true. You are buying credibility with information he was going to find anyway.

CAREFUL

Do not overcorrect into self-flagellation. Four sentences, said once, calmly, then move forward. If you spend five minutes listing your weaknesses he will conclude you don't believe in the company.

The rehearsal that actually works

Do not read the table again. Do this instead, out loud, tonight, twice:

1. **Say the four corrections** — waitlist is verbal, retention is an earlier build, quotes are paraphrases, the engine doesn't fit its weights — into an empty room, at normal volume. Time it. It should take under twenty-five seconds.
2. **Say the opening from Section 11** without looking at the page. If you reach for the page, you have not learned it.
3. **Answer "why won't this be Rewind?"** three times, differently each time, until you find the version that sounds like you rather than like this book.

The test is not whether you can recite the answers. It is whether the answers sound like they are yours. A memorised paragraph delivered perfectly sounds worse than a slightly ragged one delivered from understanding.

NOTE

If you only do one of those three, do the first. The four corrections are the thing most likely to ambush you, because they will come up early, in response to your own deck, and you will be tempted to defend a number that cannot be defended.

07 The traction conversation

You have almost no traction in the sense an investor means, and you have a surprising amount in the sense that actually predicts companies. The mistake would be to argue about which definition is right. The move is to give him the real shape of it, in a way that shows you know the difference.

What is actually true

What happened	What it is	What it is not
You spoke to 300+ people at hackathons, college tech events and in work conversations — PMs, founder's-office interns, SDEs, strategy analysts.	A real, large, qualitative signal that the problem is recognised on contact.	A waitlist. There is no list, no emails collected, no dates.
26 people used an earlier build for some days.	A genuine small-scale use test of a much earlier prototype.	Users of the current product. Nobody has run the current build on their own machine.
19 of 26 said they would use it daily once it was finished.	Stated intent, at close range, from people who had used something.	Retention. Intent is not behaviour, and this was a different build.
You showed the product at a hackathon and the room responded.	Live product-market resonance — the thing that is hardest to fake.	Adoption. Applause is not a metric.

Version one — the twenty-second answer

Use this when he asks "what's your traction?" in passing, early.

> **SAY THIS**
> "Pre-revenue by choice, and thin. About 300 conversations at hackathons and college events — PMs, SDEs, founder's-office people — and 26 who used an early build for a few days, 19 of whom said they'd use it daily once it's finished. No written list, no instrumented retention. What I do have is a working build and a person who uses it every day, which is me."

Version two — the ninety-second answer

Use this when he leans in, or when you volunteer the correction yourself.

> **SAY THIS**
> "I'll give you the honest version rather than the deck version. The deck says 300-plus waitlist and 19 of 26 retained at day 30. What's actually true: I've had 300-plus conversations, mostly at hackathons and college tech events, and the pattern is consistent — people describe the exact behaviour before I finish the sentence. There's no written list. The 19 of 26 was an earlier build, used informally for a few days by classmates and people I know, and I measured it by asking, not by instrumenting. So I treat it as a signal, not as retention. What I'd rather show you is the engine: there's a corpus from a real working day in the test suite and it fails if two works ever merge. That's evidence I can defend."

WHY THIS WORKS

That answer does something specific: it moves the conversation from a number you cannot defend to an artefact you can. He will remember the corpus, not the missing waitlist.

Version three — if he pushes hard

Some associates push. Let them. Do not get defensive and do not add new numbers to compensate.

SAY THIS

"I could give you a bigger number and you'd have no way to check it. I'd rather be the founder whose small numbers are real. If it helps: the thing I'd want you to underwrite is not the traction, it's whether the engine can be right about which work a window belongs to. That's testable and I can show you the test."

CAREFUL

Never say "we have 300 users." Never say "19 of 26 retained." The moment you compress "people I spoke to" into "users," you have made a statement that diligence will falsify, and a falsified traction claim at pre-seed is close to fatal — not legally, but reputationally, in a small industry that talks.

What to do between now and Friday

If you have two hours, the single highest-value thing you can create is **one page** titled something like Signal log . Four columns: who, where, what they said, what they'd pay. Ten rows. Real names or roles, real events, real sentences. It converts "300 conversations" from a claim into a document, and it is the first thing you'd send after the call.

WHY THIS WORKS

Ten real rows beat three hundred claimed ones. An associate who sees a messy, honest log believes it instantly, because nobody fabricates something that unglamorous.

08 The vibe-coding answer

This is the question you are most afraid of, and it is the one where honesty pays the most — because the alternative is being caught. Thomas co-authored a framework used by Meta and Microsoft. If you imply you hand-wrote 138,000 lines of Rust, one question about borrow-checker pain will expose it, and the meeting will end politely and permanently.

But here is the thing you may not have noticed: **the honest answer is strong, not weak.** It describes a real and increasingly common way to build software. What it requires is that you can demonstrate the two things that matter: that you understand what the system does, and that you can tell what is wrong with it.

The core answer — thirty seconds

> **SAY THIS**
> "I should be direct about how it was built. I'm a product-first founder, not a systems engineer. I designed how Evo behaves — the evidence model, the rule that production defines identity and attention doesn't, the fact that a shared resource abstains rather than voting — and I directed AI coding agents to implement it. What I can do is read the architecture, run the test suite, and tell you exactly which parts are solid and which aren't. And the first hire from this round is a senior Rust and macOS engineer to own the core."

Why that answer works

- It **answers** rather than deflects. Deflection is what he's listening for.
- It puts the value in the **design decisions** , which are genuinely yours and genuinely hard — "production defines identity, attention doesn't" is a real insight that most engineers would not have arrived at.
- It **names the hire** , which converts a weakness into a use-of-funds line. That is the move that turns an objection into a plan.
- It is **checkable and true** , so nothing later can contradict it.

If he asks "so can you code?"

Do not inflate. Do not grovel. Give him a calibrated answer.

> **SAY THIS**
> "Not at the level this codebase needs. I can read it, I can run it, I can tell you when something is wrong, and I can write enough to test a hypothesis. I can't own a 138,000-line Rust workspace. That's why the first hire is the person who can."

If he asks "show me the code" or "what's the worst part of it?"

This is the moment that decides whether he believes you. **Have a real answer ready** , and make it specific. The best answer is a genuine known weakness, named precisely, with the fix you already have in mind.

SAY THIS

"The worst part is the browser. Restoring a URL is easy; restoring the exact tab in the exact profile with the session intact is an adapter problem per browser, and right now the honest status is that the adapter exists and is tested against a configured endpoint, but I don't have a clean end-to-end receipt on a second machine. The second-worst is unsaved buffers — if you haven't saved, nothing can bring it back, and pretending otherwise is how you lose a user's trust on day one."

WHY THIS WORKS

That answer is worth more than a demo. It proves three things at once: you know the codebase, you know what's broken, and you don't lie about it. Associates remember the founder who told them the worst thing unprompted.

CAREFUL

Do not say "I built it myself" as a shorthand, even casually. Do not say "I'm technical" without a qualifier. Do not say "it's mostly AI but I review everything" and then fail to name a single thing you changed. The claim is not the problem — the unearned specificity is.

The reframe, if you want one

There is a legitimate strategic point here, and it is worth making once, briefly, without defensiveness:

SAY THIS

"There's a real argument that the scarce skill at this stage isn't writing the code, it's knowing what to build and being able to tell whether it works. I've got the second one — the corpus and the receipts are how I tell. What I'm buying with this round is the first one."

Then stop. Do not elaborate. The more you explain, the more it sounds like an excuse.

TECHNICAL TRUTH

09 What the code actually does — and the four things it can't

You will not be asked to explain the architecture. You may be asked to demonstrate that you understand it. This section gives you a plain-English model you can hold in your head, then the specifics for follow-ups.

The pipeline, in six sentences

1. **Capture.** A background process reads macOS Accessibility events — which window has focus, what document it holds, what URL, what path — plus file saves and input activity counts. It never takes screenshots.
2. **Ledger.** Those events go into a local append-only SQLite log with full-text search. It is the single source of truth, and it is deterministic.
3. **Episodes.** The log is cut into bounded stretches where attention stayed on a stable working set. Boundaries come from content-tested resumption first, dip absorption second, and the clock only as fallback.
4. **Bodies of work.** Episodes are assigned to works. Identity comes from production — what you typed, saved, downloaded, committed — not from what was merely open. Attention-only resources become companions and decay.
5. **Retrieval.** A request — typed or spoken — resolves to one work, or honestly reports ambiguity, or reports no match. It never silently widens scope.
6. **Restoration.** Each resource becomes an action: open a URL, open a file, or raise an existing window if one already shows that exact document. Every attempt lands in a bucket and is reported.

The four ideas that are actually yours

If he asks what is novel — and he will — these are the four, and they are real:

Idea	Plain English	Why it matters
Production defines identity	What you made decides which job a thing belongs to. What you merely looked at doesn't get a vote.	This is why fifteen open windows across four jobs don't confuse it. It is the single rule the whole engine rests on.
Shared resources abstain	If a browser or terminal genuinely serves two jobs, it votes for neither. The engine forks and shows both instead of guessing.	It converts the hardest classification case from a guess into a question. No competitor does this.
A merge requires a declaration	There is no code path from inference to merging two works. Only the person can merge.	Silent wrongness — a client's work appearing in another client's room — becomes structurally impossible.
Every restore is receipted	Opened, raised, or unopened-with-reason. Plus verified, unverified, or failed.	The product can tell you what it didn't do. That is the trust mechanism, not the privacy claim.

How you evaluate it — the answer to his favourite question

This is the one to get right. He wrote that the field's missing half is "owning evaluation." So the answer to "how do you know it works?" is the most important technical answer you will give.

> **SAY THIS**
> "Two layers, and I'll be precise about what each proves. There's a corpus from a real recorded day of mine — about ten thousand events, a hundred and sixty hand-labelled resources across six bodies of work. It replays through the engine inside the test suite and hard-fails on three properties: no thread may fuse two different labelled works, no leisure resource may anchor a work, and replay must be deterministic. The same check runs as a CLI that exits non-zero. That proves I haven't broken safety. It does not prove accuracy — it's one person, two days, and a conservative engine mints few threads, so a clean run is necessary and not sufficient. I know that."

Then, if he wants more:

> **SAY THIS**
> "The next layer is per-user evaluation. Every correction a person makes — merge these, keep these apart, this isn't mine — is a labelled pair by definition, and it's the only ground truth that scales past one person's corpus. That's what I'd build with the round."

> **CAREFUL**
> Do not say "the gate runs on every commit in CI." There is no CI configuration in the repo. The true sentence is: "It's a hard test in the workspace suite — 59 suites, about 1,270 tests, green at the last full run — and the same check exists as a CLI with an exit code, so wiring it into CI is an hour of work." That is a better answer anyway, because it shows you know the difference between a test and a pipeline.

The four things it cannot do

Volunteer these. All four, briefly, in one breath, before he finds them. Nothing buys credibility faster than a founder naming his own limits in technical language.

1. **No learning in the weights.** The engine accumulates evidence and obeys corrections. It does not fit a per-user model. Anyone who says otherwise is describing a roadmap.
2. **No browser session fidelity.** It can reopen a URL, and there's an adapter for a configured DevTools endpoint that can focus an exact tab — but authenticated sessions, form state and navigation history are not restored.
3. **No unsaved buffers.** If you didn't save it, nothing can recover it. Any product claiming otherwise is lying or screenshotting.
4. **macOS only, Apple silicon, no public release.** No signing and notarisation pipeline, no auto-update, no Windows.

If he asks about the test numbers

Suites	59 at the last full offline run. 57 at the September close-out.
Tests	1,272 passed, 0 failed, 4 ignored on 5 October. 1,087 passed, 0 failed on 30 September. Re-run before Friday and quote the fresh one.

What the tests do not prove	They do not prove restoration works on real hardware end to end. The repo's own release audit says the full capture-to-UI receipt, the cold Accessibility probe and a real browser profile receipt remain unpassed. That document exists and you should mention it — it's evidence of engineering maturity, not a liability.
Language, size	Rust. 16 crates. Roughly 138,500 lines.
Network	No HTTP client anywhere in the workspace. The daemon makes no outbound calls. That is checkable, and it is the strongest form of the privacy claim you can make.

WHY THIS WORKS

That last one is worth memorising. "There is no HTTP client in the codebase" is a much stronger sentence than "we don't send your data anywhere," because it is architecture rather than a promise.

If he asks to see the repo

Say yes immediately. Then tell him where to look, because a 138,000-line workspace with no guide reads as noise and makes you look like you're hiding in it.

Send him to	Because it proves
README.md	The product claim, the layer boundaries, and the fact that the repo openly says it is an engineering build and not a release.
docs/BUILD-LOG.md	A chronological, evidence-carrying log of every round, with test counts and known gaps. This is the document that makes people believe the project is real.
docs/architecture/LOCAL-WORK-RESTORATION.md	The audit that found the wrong-window risk, the duplicate voice authority and the downloading embedder — and fixed them. Honest engineering, written down.
crates/evo-threads/tests/ground_truth.rs	The three hard assertions. Two minutes of reading and he understands your quality bar.
crates/evo-execution/src/restore.rs	The opened / raised / unopened-with-reason contract, explained in the module doc comment.

SAY THIS

"Start with the build log and the restoration audit. If you read those two and still think it's a wrapper, I'd rather know that now."

WHY THIS WORKS

Offering your own repo as evidence, with a suggested reading order, is the opposite of what an insecure founder does. It also means his diligence lands where you are strongest: the written record, not the line count.

What to say if he finds something bad

He might. The repo contains an honest list of its own unpassed gates. If he quotes one back at you, agree with him:

SAY THIS

"Yes — that's in the release checklist, and it's accurate. The end-to-end capture-to-UI receipt isn't done, and neither is a real browser profile receipt. I wrote that list so I couldn't lie to myself about where I am. Those are the first two things the round buys."

Built, partial, and not built — the honest grid

If you remember one table from this section, remember this one. It is the entire product state in nine rows, and every cell is defensible.

Capability	State	The honest words
Observe window, document, URL, path	Built	"It reads what's focused and what it holds. No screenshots, no keystrokes."
Local ledger, FTS5, deterministic replay	Built	"Append-only, SQLite, replays identically."
Group activity into bodies of work	Built, imperfect	"It works and it's the part that needs the most work. That's what the corpus is for."
Restore files, folders, URLs	Built	"Opens, raises, or reports why not."
Voice: wake word, transcription, intent	Built, local	"Whole path on-device. No cloud speech."
Exact browser tab in exact profile	Partial	"Adapter exists against a configured endpoint. No clean second-machine receipt yet."
Terminal working directory	Partial	"Where the shell exposes it."
Per-user learned model	Not built	"It compounds memory, not parameters."
Cross-device encrypted sync	Not built	"Month two to five. It's why the round is 1.5 and not 500K."

CAREFUL

Notice that the "not built" column includes the two things investors most want: learning and sync. That is uncomfortable and it is correct. A founder who can name what he hasn't built is trusted about what he has; a founder who elides it gets a diligence call that ends the process quietly.

THE BORING THING THAT CAN END A ROUND

10 The incorporation and visa question

Sooner or later — probably in this meeting, certainly in diligence — someone asks whether you can legally run a US company and be in the United States. Most first-time founders improvise, and it shows. You will not.

The one-line answer

> **SAY THIS**
> "We're incorporated in Delaware — that's the entity you'd wire into. Neither of us needs a visa to run it or to raise. What we need a visa for is being on the ground, and the round funds that: we file I-941 under the International Entrepreneur Rule on close, because the investment clears the threshold several times over. Worst case we're on O-1s. Either way it doesn't gate the business."

Why each clause is true

Clause	The fact behind it
Incorporated in Delaware	You do not need citizenship, residency, an SSN, a US address or physical presence to form and own a Delaware C-corp. It is a filing, not an immigration event. If you are currently an Indian entity, the standard structure is a new US parent with the Indian company as a subsidiary — normal, but know it before he asks.
Neither of us needs a visa to run it	The company is a US legal person regardless of where its officers sit. Board meetings over Zoom are fine. This is corporate law, not immigration law.
I-941 on close	International Entrepreneur Parole requires: a US entity formed within five years; at least 10% ownership; a central and active role; and at least $311,071 in qualified investment from a qualified investor within the prior 18 months. A $1.5M round clears that roughly five times. NEA clears the qualified-investor test without effort. Grants up to 2.5 years, renewable once — five years total. Up to three founders per company.
Worst case O-1	Since January 2025 USCIS explicitly permits a founder's own company to petition for them, provided there's an independent director with authority. You need three of eight criteria. No cap, no lottery, three years initially, premium processing in fifteen business days.
It doesn't gate the business	Your first market is the US, not your address. You can be in front of customers on B-1 for meetings and demos while longer-term status processes.

The trap

> **CAREFUL**
> Do not tell him you are moving to San Francisco on a date. If your pitch is "we relocate after the round," you have just told an investor that your go-to-market depends on a discretionary government decision you do not control. I-941 is parole, not a visa: it is granted case by case, it cannot be appealed if denied, and it does not lead to a green card on its own. Never put that on the critical path of the pitch.

Say instead: **the US is the market before it is the address.** That is stronger, cheaper, and true. Two founders in San Francisco on $83K a month is a choice; two founders in India with a US sales presence is a strategy — and it is the reason $1.5M buys eighteen months instead of six.

If he drills

Question	Answer
"Can you work in the US right now?"	Not as a resident. B-1 for meetings and customer conversations, which is what this trip would be. The company is a Delaware C-corp and I'm its CEO regardless of where I'm sitting.
"When will you move?"	When the status lands, not before. I'm not going to promise you a date that a government agency controls. What I will commit to is being in front of US customers from month one.
"Is this a diligence problem?"	It's a documented path with a threshold we clear by roughly five times. I'd rather show you the plan than tell you it's fine.
"Why not just stay in India?"	Engineering stays. Sales doesn't. The cost base is why $1.5M buys eighteen months, and I'm not going to spend it on rent before we have revenue.
"Is the cap table clean?"	Two founders, [X]% each, no SAFEs yet, no advisor grants. Know your actual numbers before Friday. If it isn't clean, say what's messy — a messy cap table disclosed is fine; one discovered is fatal.

Four things to have done before Friday

1. **Know your incorporation status exactly.** Entity name, state, date, who is on the cap table and at what percentage. If you are not yet a Delaware C-corp, decide now whether you will be — a US parent with an Indian subsidiary is standard, and "we're forming it now" is an acceptable answer.
2. **Start the EIN if you haven't.** Without an SSN the IRS route takes **10 to 35 business days** , not two. You cannot open a US bank account or receive a wire without it.
3. **Respect the 83(b) clock.** 30 days from founder stock issuance. No extensions, no remedy. Missing it means ordinary income tax on every vesting tranche of stock you can't sell.
4. **Budget $8–15K for immigration counsel** inside the use of funds. It is a line item that signals competence and it is genuinely needed — I-941 petitions fail on thin evidence, not on thresholds.

> **NOTE**
> One thing to get advice on but not to improvise: if you run a Delaware C-corp from India, Indian tax law has a place of effective management concept that can treat a foreign company as an Indian tax resident. If he asks about the India/US structure, "we've taken advice on India-side tax residency and the entity structure reflects it" ends the topic. Not having thought about it starts a diligence problem.

11 The first four minutes

He will say something like "tell me about Evo" within the first two minutes. That is your only open invitation. Most founders spend it on the problem, because the problem is the easy part. Do not. Spend it on the sentence he needs to write in his memo.

The opening — say this, roughly

SAY THIS

"Thanks for making the time. Thirty seconds of context, then I'll show you rather than tell you.

Every knowledge worker has a second job they don't get paid for: reconstructing what they were doing before they were interrupted. The files, the tabs, the decisions. Nobody has solved it because it isn't a search problem — the answer isn't on the web, it's on your laptop, in the state of your machine.

Evo watches work happen locally, groups it into bodies of work, and rebuilds the whole workspace when you say continue. It's a native Rust system, sixteen crates, about a hundred and thirty-eight thousand lines. Nothing leaves the device — there's no HTTP client in it at all.

The thing I'd want you to judge me on isn't the demo. It's that the engine has to be honest. Every restore reports what it opened, what it brought forward, and what it couldn't open. And there's a corpus from a real working day in the test suite that fails the build if two separate bodies of work ever merge.

Can I show you? It's about three minutes."

Why this opening, specifically

Sentence	What it is doing
"Thirty seconds of context, then I'll show you"	Sets a contract. Signals you will not monologue, and gives him permission to interrupt — which makes him relax.
"a second job they don't get paid for"	A concrete, human framing of the problem. Not "context switching costs the economy $X billion."
"it isn't a search problem"	This is the positioning move, and it is the most important sentence in the opening. It distinguishes you from Recall, Rewind, Granola, Raycast and ChatGPT in one clause.
"native Rust… sixteen crates… no HTTP client"	Three specifics that prove there is a real system. An associate's instinct is "is this a wrapper?" — this answers it in eight seconds.
"the thing I'd want you to judge me on isn't the demo"	Reframes the meeting from entertainment to underwriting. It tells him you know what he's actually doing.
"Can I show you?"	Asking permission rather than launching. He says yes 90% of the time, and now he is participating instead of evaluating.

CAREFUL

Do not open with the market size. Do not open with "we're the ChatGPT of X." Do not open with your background. Do not open with the traction slide. Every one of those is a reason for him to categorise you as generic before you have said anything that only you could say.

If he asks you to update the deck

He might say the deck overstates something, or that it reads as consumer rather than infrastructure. Do not defend it. Agree, and ask what he would change — then actually change it and send it back within 48 hours. A revised deck that reflects his feedback is the cheapest possible signal that you listen.

> **SAY THIS**
> "You're right, and I'd rather fix it than defend it. If you were writing that slide, what would it say?"

This is a small thing that lands disproportionately, because most founders treat their deck as an artefact to be protected rather than a document to be improved. The ones who come back with a better version get remembered.

If he opens differently

Associates often skip the invitation and go straight to a question. Handle these three, then take the wheel.

He says	You say
"I read the deck — give me the short version."	The opening above, but cut the middle. Keep the second-job framing, the "not a search problem" line, and the honesty line. Drop the crate count.
"How did Lila end up forwarding this?"	Be brief and don't gloat. "I read her Neolab post and it described the problem I'd been building against for a year, so I wrote to her about that specifically rather than sending a deck cold." Then move on.
"What does it actually do?"	Give one concrete example, not a description. "I was working on the investor deck Friday. I come in Monday, say 'continue the pitch deck,' and the deck, the two tabs I was reading and the folder are there. That's it. That's the whole product."

12 The demo

You have a working build. That is your single biggest advantage over every other pre-seed founder he will meet this week. Do not waste it on a tour of the interface. A demo is an argument, not a walkthrough — every click should prove a claim you could otherwise only assert.

The four beats, in order

Beat	What you do	What it proves
1. Resume	"I was on the pitch deck Friday." Say: "Hey Evo, continue the pitch deck." Let it rebuild. Point at the clock when it lands.	That restoration is an action, not a search result. This is the whole product in eight seconds.
2. Memory across time	"It knows this from a session two days ago. I didn't tell it."	That the state is durable rather than session-scoped. This is the "longitudinal" word he uses.
3. Separation	Show two works that share a browser. "Same browser, both jobs. It doesn't guess which — it abstains and asks."	The hardest technical claim, made visible. No competitor demonstrates this because none of them can.
4. Correction	Drag a window into the wrong pod. It merges, and stays merged. "One gesture. Permanent."	That wrongness is cheap and reversible, and that the correction outranks inference. This is the trust answer, demonstrated.

Before the call — the thirty-minute ritual

1. Confirm the daemon has events for today. A cold machine demos nothing.
2. **Create production in each of the three works you plan to show.** Type a line, save a file, run a command or make a commit. Production is what mints identity — a work you only looked at might not exist as a work.
3. Open the pod bar and check that the three works are there. **If one is wrong, fix it now with a merge** — and then that becomes a demo beat rather than a liability.
4. Close everything you don't need. A cluttered screen reads as a cluttered product.
5. Turn off notifications, Slack, mail, and every badge. Put the machine on charge.
6. Have the 90-second recording open in a tab, and say so out loud at the start.

Say this before you share your screen

SAY THIS

"Two things before I share. I've recorded a ninety-second version in case the screen share fights me — live demos fail for boring reasons and I'd rather not spend your thirty minutes on a permission dialog. And if it does something wrong while you're watching, tell me, because that's the most useful thing that could happen in this call."

WHY THIS WORKS

Both halves of that are strategic. The recording is insurance, and announcing it makes you look prepared rather than defensive. Inviting him to catch errors converts the demo from a performance into a shared diagnosis — and it means a glitch costs you nothing.

When it breaks

What happens	What you say	What you do not do
Restore opens the wrong work	"That's the interesting failure. It's merged these two — let me show you the correction, because that's the actual design."	Do not apologise repeatedly. Do not close the window and pretend.
Nothing opens	"It's not finding it. Give me fifteen seconds — I'll switch to the recording rather than debug live."	Do not troubleshoot for two minutes in silence. The silence is what he remembers.
Permissions dialog appears	Handle it calmly, narrate what it's asking for and why. This is actually a useful beat: it shows the permission model.	Do not panic-click. Do not say "this never happens."
Voice doesn't hear you	Type it. "It's the mic, not the engine." Then move on.	Do not repeat the wake word five times. Do not blame the laptop.

The YouTube-tab question

You mentioned wanting to show a video you had been playing. Be careful here, and the reason is architectural rather than presentational. **The engine's own test suite hard-asserts that a leisure resource never anchors a work.** If you play a random video and Evo files it into a work, you are demonstrating the exact failure your own gate forbids.

So either leave it out, or use it deliberately — and if you use it, frame it:

> **SAY THIS**
> "This tab is reference material for the deck, so it rides along as context. But notice it can't found a work — a video I just watched doesn't get to claim that I was working on something. Only production does that."

> **WHY THIS WORKS**
> That is a superb beat, because it turns a potential embarrassment into an explanation of the core rule. But only do it if the tab is genuinely attached to a work. If it's leisure, leave it out entirely — the risk outweighs the ten seconds of interest.

13 The forty questions he will actually ask

Grouped the way he will ask them. Every answer is thirty to forty-five spoken seconds. The first sentence is the answer; the rest is evidence; then stop.

A. The opening five

1. "Tell me about Evo."

Section 11. Do not improvise a new version on the day.

2. "Who is this for, concretely?"

> **SAY THIS**
> "Developers, designers and founders who run three or four projects in parallel on one machine. The tell is that they already have a ritual for recovering — a saved tab group, a notes file called 'where I was', a rule about never closing the laptop. People who feel the tax daily build their own bad version of this."

3. "Why you?"

> **SAY THIS**
> "Because I lived it for two years and then spent a year building the unglamorous half. The easy part is reopening windows. The hard part is deciding which of your fifteen open things belong to the same job, and I've got a rule for that which most engineers wouldn't have reached: what you produced defines the work, what you merely looked at doesn't."

4. "Why now?"

> **SAY THIS**
> "Three things landed in about two years. On-device models got small enough to run on a laptop without a battery cost worth mentioning. macOS exposed the capture APIs that make this possible without screenshotting. And an OpenAI product made the category legible to normal people — which is good for us, because they showed everyone they want this and they did it in the cloud with a wall around it."

5. "Why is this a company and not a feature?"

> **SAY THIS**
> "Because the hard part isn't a capability, it's a corpus. Knowing which of your windows are one job takes months of your own evidence, on your machine. A platform vendor can ship the button; they can't ship your six months of context. And the switching cost is tenure, not features — after a month, starting over means amnesia."

B. Product

6. "What exactly gets restored?"

> **SAY THIS**
> "Files, folders, URLs in the browser they were seen in, and window positions for apps I have adapters for. Five things at most by design — it reopens the neighbourhood, not the museum. What it does not restore is unsaved buffers, authenticated session state, or form contents."

7. "Why not Spaces, Stage Manager, browser session restore, Raycast?"

> **SAY THIS**
> "Because all of them save state and none of them save meaning. Spaces saves a desktop, not a project. Session restore gives you back forty tabs with no idea which ten mattered. Raycast finds things when you know what you're looking for. The gap is knowing which of your things are the same job — and that needs evidence over time, not a launcher."

8. "What's the activation moment?"

> **SAY THIS**
> "The first time they say continue and the thing comes back. That's it. Everything before that is setup and everything after is habit."

9. "How long is setup? Which permissions?"

> **SAY THIS**
> "Accessibility, which is the big one, plus optionally microphone for voice. Five minutes and one restart. The Accessibility grant is the real friction and I'm not going to pretend otherwise."

10. "Does it work offline?"

> **SAY THIS**
> "It works only offline. There's no server to be online to. The embedding model for voice runs locally."

C. Technical

11. "Why Rust?"

> **SAY THIS**
> "Because it's a long-running local process touching the filesystem, the window server and SQLite at once, and I wanted memory safety without a garbage-collection pause in the middle of a restore. The honest second reason is that a daemon has to outlive the UI, and Rust makes that boundary clean."

12. "Walk me through the pipeline."

Section 09. Six sentences. Do not go deeper unless asked.

13. "How do you evaluate classification? What's your ground truth?"

Section 09. This is the most important technical answer in the meeting — it is his stated area of interest. Do not rush it.

14. "What happens when it gets it wrong?"

SAY THIS

"It gets things wrong. The design question is what happens next, and the failures aren't symmetric. Fusing two works is expensive — a client's material shows up in another client's room. Splitting one work is cheap. So the engine prefers splitting: when two candidates are close it forks and shows you both, and a resource that genuinely serves two works votes for neither. A merge only happens if the person declares it — there is no code path from inference to a merge. So wrongness is visible, cheap, reversible, and one gesture from permanent."

WHY THIS WORKS

This answer is far stronger than "it's smart enough to get it right," which is unfalsifiable and invites "what's your error rate?" It is also the answer that tells a technical investor you understand failure modes — which is what separates an engineer from an enthusiast.

15. "How does it handle Electron apps, browser tabs, terminals?"

SAY THIS

"Three different problems with three different answers. Browser tabs are the furthest along — there's an adapter that can focus an exact tab in an exact profile. Terminals restore the working directory where the shell exposes it. Electron apps are the weak spot, because they expose wildly different automation surfaces; some give you a document, some give you a window title and nothing else. That's why it's a capability tier per app rather than a promise."

16. "Is there a model? Which one?"

SAY THIS

"There's a local embedding model — MiniLM, 384 dimensions — used for matching a request to a work. The classification itself is deterministic evidence scoring, not a neural network. That's deliberate: a three-billion-parameter model running continuously on a laptop is a battery and thermal problem, and it would make the product worse."

17. "Does it learn?"

SAY THIS

"Today it compounds memory, not parameters. Every correction is permanent and outranks anything the engine infers, but it doesn't fit its own weights yet. The next layer is exactly that: the corrections are labelled pairs — merge these, keep these apart — and a few hundred weights fitted on-device would resolve the near-ties the rules currently fork on. Deterministic core for safety, learned layer for arbitration."

CAREFUL

Do not say "it learns automatically and gets better." He will ask what the update rule is, what the label is, and where the loss is — and there are no answers in the code. "It accumulates evidence and obeys corrections" is true and defensible.

18. "Multi-device sync?"

SAY THIS

"Not built. It's month two to five, and it's why the round is $1.5M and not $500K. The design constraint is that the server never sees plaintext — client-side encryption, opaque blobs, and the append-only ledger means conflicts resolve by merging evidence rather than last-write-wins."

19. "Windows?"

SAY THIS

"Scoped, not built. Capture is tractable — UI Automation and Win32 give you window and document identity. Restoration is the expensive half, because every app exposes a different surface. Capture first in the six weeks after the round; full restore is months per app family."

20. "How big is the codebase, and who maintains it?"

Section 08. Answer it plainly, name the first hire, then stop.

21. "What breaks on a macOS update?"

SAY THIS

"Everything that touches Accessibility. It's an unsupported surface in practice, and Apple changes behaviour between minor versions. That's why the capture layer is isolated behind one boundary and why I test against real windows rather than mocks."

D. Privacy and trust

22. "You're asking for Accessibility. That's what keyloggers use."

SAY THIS

"Four distinctions. We use Accessibility to read which window is focused and what document it holds — not Input Monitoring, which is the permission that intercepts keystrokes, and we never ask for it. We don't capture screenshots or audio by default. We store metadata about activity, not content. And there's no server in the loop at all — there's no HTTP client in the codebase. It has to be defensible to a security researcher, not just to a user, because one bad write-up ends a local-first company."

23. "What if malware reads the local store?"

SAY THIS

"Then it has your activity metadata, which is a real risk and I'm not going to wave it away. Today the protection is FileVault and the sandbox. Per-record encryption with a Keychain key is a requirement before sync ships, and it's on the plan for month two."

24. "GDPR? Enterprise?"

SAY THIS

"Consumer first, so GDPR is a later problem — but the architecture helps, because the data never leaves the machine. The hard part will be the employee's activity being visible to an employer, and I'd want to solve that with the employee holding the key, not the company."

25. "How do you prove nothing leaves the device?"

SAY THIS

"Three ways, in increasing strength. There's no HTTP client in the dependency graph, so it's checkable by reading the manifest. The daemon has no network code path. And the release checklist requires offline packaging tests with the network denied. I haven't shipped a third-party audit yet — that's in the use of funds."

E. Competition and market

26. "OpenAI shipped Computer History in August. Why are you not dead?"

SAY THIS

"Because they answered a different question. Computer History makes ChatGPT remember — it sends your interaction events to OpenAI, turns them into memories, and gives you a timeline inside ChatGPT. Evo restores your actual workspace, whichever AI you use, and the events never leave the machine. If OpenAI ships full workspace restore on-device, I have a problem — and the thing that would save me is tenure, not features."

WHY THIS WORKS

Naming the risk in the answer, rather than waiting to be caught, is what makes the preceding sentence believable. Associates discount optimism; they trust calibration.

27. "What about Microsoft Recall?"

SAY THIS

"Recall is the most technically impressive thing in the space and it's Windows-only, on new hardware, with a privacy history it will never fully outrun. There is no trusted equivalent on the Mac. That's the gap."

28. "Why won't this be Rewind?"

SAY THIS

"That's the right question and I think about it a lot. Rewind recorded everything and gave you search — it was a very expensive archive you had to remember to consult. This is the opposite: it does one thing at the moment of return, and the archive is a side effect. Rewind also needed constant screen capture, which is why it was heavy and why it scared people. We capture metadata, which is cheap and defensible."

29. "What does Apple do?"

> **SAY THIS**
> "Apple will do the on-device half well and will never do the cross-vendor half. They won't restore your Chrome profile into Arc, and they won't build a layer that works better because you also use ChatGPT and Claude. Platform owners make the floor; the layer above the floor is where products live."

30. "Bottom-up market size?"

> **SAY THIS**
> "6.1 million US computer and mathematical professionals is the beachhead. I need 1.6% of that to reach a hundred thousand users. At 9% paying $108 a year, that's about $970K of ARR. Then 70.7 million management and professional occupations is the same product one segment over."

31. "Is this big enough?"

The most important answer in the meeting. Slow down for this one.

> **SAY THIS**
> "As a paid Mac utility, honestly, no — that's a five to forty million dollar ARR business and I'd be fine saying so. The reason I think it's bigger is that the thing we accumulate isn't a feature, it's the boundary map of a person's work. Every agent shipping right now has memory and no jurisdiction — it knows facts about you, it doesn't know that this repo is one client and that deck is a different job. So it can't safely act on your desktop. We produce that boundary as a side effect of being useful every morning, and then agents query it. That's the layer."

> **CAREFUL**
> Do not answer this with TAM. He has heard "$1B knowledge workers" four times this week. Answer it with the sequence: wedge, then asset, then layer — and concede the ceiling of the wedge, because conceding it is what makes the layer claim credible.

F. Traction, team, fundraising

32. "What's your traction?"

Section 07, version one or two depending on the temperature.

33. "You're a solo founder?"

> **SAY THIS**
> "Two of us — I'm CEO and CTO, Manavv runs commercial. But on the engineering side I'm effectively solo today, and the honest risk is bus factor. The first hire from the round is a senior Rust and macOS engineer, and I've already scoped what they own."

34. "What are you weakest at?"

> **SAY THIS**
> "Systems engineering at depth, and distribution. I can design the behaviour and tell when it's wrong; I can't own a large Rust workspace. And I've never run a paid acquisition channel. Those are the two hires."

35. "How do you recruit a senior Rust engineer?"

SAY THIS

"The honest answer is that it's the hardest part of the plan and I'd take help. What I have going for me is that it's an unusually good problem — local-first systems, macOS internals, deterministic state. That's a specialist's dream project, not a CRUD job."

36. "Why $1.5M?"

SAY THIS

"Eighteen months to a seed on 100K users and roughly a million of ARR run-rate. The 1.5 is 83K a month — two founders, two engineers from month four, and the launch budget. It's a milestone budget, not a wish."

37. "At $12M post?"

Section 17. Have the number, and have the flexibility.

38. "Who else is in the round?"

SAY THIS

Answer honestly. If it's nobody: "Nobody committed yet. I started with the people whose thesis this actually is rather than running a broad process." That is a legitimate strategy and it flatters him slightly without lying.

39. "What do you want from us beyond money?"

SAY THIS

"Two things. A partner meeting, because I think the category argument is better made to someone who owns a fund thesis than to me in thirty minutes. And help on the first senior hire — you've backed infrastructure companies, you know what a good systems engineer looks like and I don't have that network in the US."

40. "What would make this a no?"

SAY THIS

"If the honest read is that it's a nice utility with no path to a layer, that's a fair no. What would change my mind about that being the read is if you think the context boundary is a real asset or just a feature of this app. I'd genuinely like to know which one you think it is."

WHY THIS WORKS

Asking for the objection at the end is disarming and it works. It also gets you information you cannot otherwise obtain — associates rarely volunteer the real reason.

HARDEST FIRST

14 The ten questions that decide it

These are the ones where a weak answer loses the meeting regardless of everything else. If you only rehearse ten answers, rehearse these. Say each one out loud until it takes under forty seconds.

1. **"Is this a feature OpenAI ships next quarter?"** — Section 13.26, plus the concession. The concession is the answer.
2. **"How do you evaluate it?"** — Section 09. The corpus, the three hard assertions, and what they don't prove.
3. **"Does it actually learn?"** — "It compounds memory, not parameters." Never bluff this one.
4. **"Who built the code?"** — Section 08. Direct, unapologetic, names the hire.
5. **"What's your traction, really?"** — Section 07. Volunteer the correction.
6. **"Why won't this be Rewind?"** — The archive-versus-return distinction.
7. **"Is it big enough?"** — Wedge, asset, layer. Concede the wedge ceiling.
8. **"Why you?"** — The production-defines-identity insight, which is genuinely yours.
9. **"What happens when it's wrong?"** — Fork, abstain, declare, receipt.
10. **"Can I speak to a user?"** — Have one name ready, with permission, who will take the call. If you don't have one, say: "Not yet — the current build hasn't gone out. That's the first thing the round changes."

> **CAREFUL**
> Number ten is the one founders forget to prepare, and it is the one that converts a good meeting into diligence. Have a name. Ask that person tonight whether you can put them forward.

THE MOST USEFUL SKILL IN THE ROOM

15 When you don't know

He can ask anything and you cannot know everything. That is normal and he expects it. What he is testing is not your knowledge — it is whether you **know the edges of your own knowledge** , because that is the trait that predicts whether you'll be honest with him for the next five years.

The four beats — fifteen seconds, never more

1. **Acknowledge.** "I don't know — that's a real gap." No flinch, no filler.
2. **Give the nearest true thing.** "What I do know is ___."
3. **Say how you'd find out, and when.** "The way I'd answer it is ___; I can have that to you by Tuesday."
4. **Turn it.** "How would you test it?" or "What would you need to see?"

> **SAY THIS**
> "I don't know. That's a gap, not a hedge. What I do know is the shape of the answer — it's a storage-growth question and I'd measure it on my own machine over a week rather than guess. I can send you the number by Wednesday. How would you sanity-check it?"

> **WHY THIS WORKS**
> That answer costs you nothing and gains you three things: honesty, a deadline you set yourself, and a question that makes him talk. Founders who admit a gap early get believed on everything else.

For design questions — show the axes

If it's not a fact you're missing but a decision you haven't made, do not fake an answer. Show the trade-off:

> **SAY THIS**
> "That's a decision I haven't made. The trade-off is privacy against usefulness: cross-device sync is the thing people ask for, and it's also the first time anything leaves the machine. My instinct is encrypted sync where the key never touches the server, but I haven't pressure-tested it. How have you seen other local-first companies handle that?"

The three you genuinely don't know, pre-loaded

Question	Your answer
"What's your storage growth per user per day?"	"Order of megabytes — the ledger is event metadata, not content. I haven't measured it precisely and I'd rather instrument it during the beta than give you a number I invented."
"What does Windows cost you?"	"Scoped, not built. Capture is weeks; restoration is months per app family. I haven't written a line of Windows code."
"What's your CAC?"	"Unknown. I haven't spent a dollar on acquisition. My hypothesis is that the demo clip carries the launch and the first cohort is organic, and that's exactly what the launch budget is for — to find out."

What never to do

- **Never invent a number.** Not a percentage, not a growth rate, not a market size. One fabricated figure, discovered later, retroactively invalidates everything true you said.
- **Never say "I'll get back to you" without a date.** A promise without a deadline is a deflection, and he will notice.
- **Never argue with a question.** If he says "this looks like a feature," do not say "no it isn't." Say "that's the risk I think about most" and then give him the counter-argument.
- **Never blame the market, the timing, or other investors.** It is the fastest way to look junior.

Nine things to say when you have nothing

Sometimes the honest answer is short. These are complete answers, not deflections — each one names a limit and then leaves the door open. Learn the shape, not the words.

The question	The fifteen-second answer
"What's your conversion rate?"	"Not measured. No Pro tier yet — that's month three. I'd rather tell you it's unknown than give you a number from a survey."
"What's your churn?"	"Unknown, and I'd be making it up. The honest proxy is that the 19 people who said they'd use it daily are the ones I'd instrument first."
"What's your CAC?"	"Zero spent so far. The launch budget exists to find out, and I don't have a number I'd defend."
"How many people are on the team?"	"Two founders. No employees. The first hire is in the plan and named in the use of funds."
"What's your legal structure?"	"Delaware C-corp, [state the cap table]. If there's anything messy I'd rather show you now than have it surface in diligence."
"Who are your advisors?"	"Nobody formal. I've taken advice from [name the real people] and I'd rather have a formal advisor who owns something specific than a logo."
"Why hasn't anyone built this?"	"People have tried — Rewind, Recall, and OpenAI just shipped a version. The difference is doing it on-device and doing the grouping rather than the search."
"What's your unfair advantage?"	"Time on the problem, and a rule about identity that took a year to get right. Neither is a moat on its own — the moat is the evidence a user accumulates."
"What keeps you up at night?"	"That I'm one person on the engineering side of a system that needs three. That's the hire, and it's the first line of the budget."

WHY THIS WORKS

Every one of these has the same structure: a plain admission, a reason it's unknown, and no apology. That is what calibrated sounds like. The alternative — a confident number that turns out to be invented — is the only unrecoverable mistake available to you in this meeting.

CAREFUL

One caution on "who are your advisors": if you have named someone as an advisor, even informally, tell Thomas the same thing you'd tell a lawyer. Verbal advisor equity promises are the single most common cap-table surprise at pre-seed, and finding one in diligence after being told "nobody formal" reads as concealment rather than as an oversight.

Nine things to say when you have nothing

Sometimes the honest answer is short. These are complete answers, not deflections — each one names a limit and then leaves the door open. Learn the shape, not the words.

The question	The fifteen-second answer
"What's your conversion rate?"	"Not measured. No Pro tier yet — that's month three. I'd rather tell you it's unknown than give you a number from a survey."
"What's your churn?"	"Unknown, and I'd be making it up. The honest proxy is that the 19 people who said they'd use it daily are the ones I'd instrument first."
"What's your CAC?"	"Zero spent so far. The launch budget exists to find out, and I don't have a number I'd defend."
"How many people are on the team?"	"Two founders. No employees. The first hire is in the plan and named in the use of funds."
"What's your legal structure?"	"Delaware C-corp, [state the cap table]. If there's anything messy I'd rather show you now than have it surface in diligence."
"Who are your advisors?"	"Nobody formal. I've taken advice from [name the real people] and I'd rather have a formal advisor who owns something specific than a logo."
"Why hasn't anyone built this?"	"People have tried — Rewind, Recall, and OpenAI just shipped a version. The difference is doing it on-device and doing the grouping rather than the search."
"What's your unfair advantage?"	"Time on the problem, and a rule about identity that took a year to get right. Neither is a moat on its own — the moat is the evidence a user accumulates."
"What keeps you up at night?"	"That I'm one person on the engineering side of a system that needs three. That's the hire, and it's the first line of the budget."

WHY THIS WORKS

Every one of these has the same structure: a plain admission, a reason it's unknown, and no apology. That is what calibrated sounds like. The alternative — a confident number that turns out to be invented — is the only unrecoverable mistake available to you in this meeting.

CAREFUL

One caution on "who are your advisors": if you have named someone as an advisor, even informally, tell Thomas the same thing you'd tell a lawyer. Verbal advisor equity promises are the single most common cap-table surprise at pre-seed, and finding one in diligence after being told "nobody formal" reads as concealment rather than as an oversight.

Write them down

Keep one sheet of paper next to you during the call. Every question you can't answer goes on it. **Then answer all of them in writing within 24 hours.** A founder who returns with a written answer the next day is more impressive than one who improvised in the moment, because it proves the loop works.

EIGHT, RANKED

16 Your questions for him

You get roughly four minutes at the end. Most founders waste them on questions whose answers are on the website. These eight are chosen because each one either gives you information you cannot get elsewhere, or makes him say something he'll have to stand behind later.

The three you must ask

1. "What happens next, and on what timeline?"

Ask it every time, at the end, regardless of how the call went. It is not pushy — it is the question a founder who has done this before asks. Whatever he says, write it down and repeat it back: "So I'll send the technical write-up Tuesday and we'll speak again the week after — does that sound right?" Repeating it back converts a vague intention into a commitment.

2. "What would make this a yes for you?"

This is the highest-value question in the meeting. It tells you what to build or prove next, and it tells you whether he has a specific bar or is just being polite. If the answer is vague, that is information too.

3. "What's your biggest concern right now?"

Almost nobody asks this and it works almost every time. It gives you the objection while you still have time to answer it — and if you can answer it well in thirty seconds, you have removed the thing that would have killed you in the memo.

The three that tell you how to play the next month

4. "How does NEA decide on a round this size, and who's in the room?"

A $1.5M pre-seed at a $38B firm is an unusual shape. You genuinely need to know whether this goes to partnership, whether an associate can champion a cheque this small, and how long it takes. Asking directly is not rude; pretending you already know is worse.

5. "Would it help if I sent you the technical write-up, or would you rather dig into the code yourself?"

Both options flatter him, and his answer tells you how technical his diligence will be. If he wants the code, you have a real advantage — the repo is genuinely well-documented and the audit document is honest.

6. "Are there parts of this you'd want me to go deeper on before we speak again?"

This converts the meeting into a homework assignment. It also makes the second call easier to earn, because you arrive having answered the exact questions he cared about.

The two to use sparingly

7. "What did you think of Computer History when it launched?"

A genuinely interesting question for a person who has written about exactly this. It also invites him to make a judgment in front of you, which is useful — you learn his taste. Use it if the call is going well and you have spare time. Do not use it if you are behind.

8. "Would it make sense to meet a partner?"

Only ask this if the call has gone well and you are near the end. Phrased as "would it make sense," it is not a demand. If the answer is "too early," you have lost nothing. If the answer is "yes, let me think about who," you have just skipped a step.

CAREFUL

Do not ask: "What's your check size?", "What's your process?", "Do you lead?", or "How much runway do you have?" — all four are answered on NEA's website or make you sound like you're shopping. And never ask him to sign an NDA. It is the fastest way to end a venture conversation.

What his answers tell you

Ask the questions, then listen to the shape of the reply. The shape is the information — not the content.

If he says	It means	What you do next
"Let me think about it and come back to you."	Polite non-commitment. He has not decided.	Send the forwardable paragraph. Make his thinking easy.
"Send me X and we'll talk next week."	A real next step. He wants a reason to say yes again.	Deliver X within 48 hours, better than he expects, and reference his exact wording.
"I'd want a partner to meet you."	You have cleared the screen.	Ask who, and what they care about, and prepare for that person specifically.
"Honestly, this might be too early for us."	The most likely real outcome. Not a judgement on the company.	Ask what specific milestone would change it. Then go hit it and come back.
"Who else are you talking to?"	He is checking whether there is a process, and whether anyone has validated you.	Tell the truth. If it's nobody, say nobody, and say why that's deliberate.

WHY THIS WORKS

The single most useful thing you can extract from the last four minutes is a specific answer to "what would make this a yes." If you get one, you have a to-do list instead of a hope. If you get a vague answer, that tells you he isn't convinced yet — and you should treat the meeting as an opening rather than a near-close.

NUMBERS

17 The ask, the cap, and the negotiation

You are asking for **$1.5M on a $12M post-money SAFE, 18 months of runway.** That is a reasonable, defensible ask. What follows is what he will think when you say it, and how to handle each thought.

The arithmetic he will do in his head

He thinks	Because	What you should have ready
"1.5 divided by 18 is 83K a month."	It's the first thing anyone does with a runway number.	A bottom-up budget, not a division. Section 17, next page.
"That's 12.5% dilution."	1.5 ÷ 12 = 12.5%.	That is a low dilution for pre-seed, which is good for you and he knows it. Don't apologise for the cap.
"$12M is below median for a $1M–$2.4M round."	Carta's Q2 2026 data puts the median post-money cap at $18M for rounds of $1M–$2.4M. The older $15M figure is now out of date.	Which means your cap is conservative, not aggressive. If he pushes back on valuation, you are on the wrong side of the argument — you are cheap, and you should know that before he tells you.
"This is small for NEA 19."	It's 0.027% of a $5.5B fund.	Section 04. Do not raise it; be ready for it.

The budget, bottom-up

Do not say "$83K a month." Say what it buys. This table is the answer to "how did you arrive at 1.5?"

Line	Monthly	Notes
Two founder salaries	$16.7K	2 × $100K a year. Modest, and say so.
Two engineers from month 4	$25.0K	The first hire is the senior Rust/macOS engineer. This is the line that matters.
Payroll tax and benefits	$7.5K	~18% on salaries.
Co-working / desk	$2.5K	Cheap. Do not claim a San Francisco office.
Legal and accounting	$2.5K	Includes the immigration counsel.
Infrastructure and dev accounts	$1.3K	Low, because there is no inference bill — the product runs on the user's machine.
Recruiting	$3.0K	Amortised. A senior systems hire costs real money to source.
Travel and events	$3.0K	US customer and investor travel from India.
Contract design and security review	$5.0K	The audit is in the plan.
Insurance and buffer	$4.0K	Don't skip this line; it signals maturity.
Base total	~$70K	21 months at this rate.
With buffer and FX	~$83K	18 months. This is the number in the deck.

SAY THIS

"Seventy thousand a month at full team, eighty-three with a buffer for FX and overruns. That's twenty-one months at base, eighteen at full. The biggest line is the second engineer, and that's deliberate — it's the hire that de-risks everything else."

WHY THIS WORKS

This is a much better answer than the division, because it shows you have thought about what the money does rather than what it costs. It also quietly explains why you're in India, without you having to defend it.

If he says the round is too small for NEA

This is the single most likely objection, and it is the one most founders handle badly by either getting defensive or immediately offering a bigger round. Neither works. Do this instead.

SAY THIS

"I know. It's a small cheque for a fund your size and I'm not going to pretend otherwise. Two honest options. One: you take a small position now because you think the category is real and you'd rather be in it before the seed — in which case I'd rather have you at $250K with your name on it than a larger cheque from someone who doesn't care about the space. Two: if the answer is that it's too small, tell me now and I'll come back at the seed with the numbers, and I'd rather have that relationship than a polite maybe."

WHY THIS WORKS

Offering the small cheque yourself is counter-intuitive and it is the strongest move available. It removes the awkwardness, it makes the conversation about conviction rather than size, and it is the only version of this conversation in which a $38B firm can say yes quickly.

If he asks "would you take more?"

Say yes, with the arithmetic already done. A larger round at the same cap means more dilution; a larger round at a higher cap means less. Know both.

If NEA wants to put in	At $12M post	At $15M post	What it buys
$1.5M	12.5%	10.0%	18 months — the plan as written
$2.0M	16.7%	13.3%	~24 months, or the same 18 with more headroom
$3.0M	25.0%	20.0%	~30 months and a proper seed-round team

SAY THIS

"I'd take two at the same cap, and I'd tell you honestly that it changes the plan — it buys a second engineer sooner and it means I'm not raising again in eighteen months. At three I'd want to talk about whether the cap moves, because twenty-five percent at pre-seed is a lot to give up before a seed."

CAREFUL

Do not negotiate against yourself. If he hasn't asked for a bigger round, do not offer one. And never move the cap in the meeting — "I'd want to think about that" is a complete sentence.

If he pushes on the cap

He probably will, not because $12M is high but because that is what associates do — it is their job to test whether a founder has a reason for his number or just picked one.

He says	You say
"That feels rich for pre-revenue."	"It might be. It's below the median cap for a round this size — Carta had it around $18M for $1–2.5M rounds last quarter — so I'm not asking for a premium. What I'm buying with it is eighteen months and a hire, and if that's not worth twelve to you I'd rather understand why than argue about the number."
"Would you come down?"	"At a smaller cheque, yes. I'd rather have a smaller round at a fair cap than a bigger one at a number I'd have to defend to a seed investor in a year. What did you have in mind?"
"We'd want a discount for being first money in."	"I understand the instinct. The counter is that first money in is also the money that makes the seed possible, and the seed is where you make the return. I'm open to structure — I'm not open to a cap that makes the next round impossible."

WHY THIS WORKS

Note what is happening in all three: you concede the possibility without conceding the number, you ask a question back, and you keep the conversation about the plan rather than the price. A founder who defends a valuation looks like a founder who needs the money. A founder who asks what the other person had in mind looks like one who can walk away.

NOTE

One structural fact worth having ready: a post-money SAFE fixes the cap, so the only variables are cheque size and cap. There is no "structure" to negotiate beyond those two, which means this conversation is shorter than it feels. If he wants more than 20% of the company at pre-seed, the cap has to move — that is arithmetic, not stubbornness.

If he asks about a lead

SAY THIS

"There's no lead yet. I started with the people whose thesis this actually is rather than running a broad process, which is either smart or naive and I'll know in a month. If NEA wants to lead, that changes the shape of the round and I'd want to talk about it properly."

AFTER THE CALL

18 Two hours, two days, two weeks

The meeting is not the thing. The meeting produces a memo, and the memo is written from memory plus whatever you send. What you send, and how fast, is therefore part of the meeting.

Within two hours

One short email. Not a deck — he has the deck. Three things, in this order:

1. **One line of thanks, no grovelling.** "Thanks for the time — I enjoyed the question about X."
2. **The three-line recap.** What you agreed the product is, what you agreed the open question is, and what you said you'd send.
3. **The links, inline.** The 90-second recording, and the one-page claim note if you wrote one. Never attach more than one file.

The forwardable paragraph

This is the highest-value artefact you can produce, and almost no founder does it. Write **a paragraph he can paste into the memo.** Not a pitch — a draft of his summary, in his voice, with the risk already in it. Founders who do this get forwarded, because they have done the associate's hardest writing task for him.

> **SAY THIS**
>
> "Pre-seed, $1.5M at $12M post. Local-first work memory for macOS: observes app/file/tab activity on-device, models bodies of work, reconstructs the workspace on command. 16-crate Rust system, ~138K lines, 1,270 tests, no HTTP client. Unusual for pre-seed: a labelled corpus from a real working day with hard regression gates against work fusion, plus per-target restore receipts. Traction is early and verbal. Risks: single founder on the engineering side, prosumer ACV, competing with OpenAI's Computer History. Angle: the context/jurisdiction layer for agents — one layer below our Namespace/Factory thesis."

Send it as text in the email body, not as an attachment, and label it something like "in case it's useful for your notes." That framing makes it a gift rather than an instruction.

Within 48 hours

- **Answer every question you wrote down** , in writing, in one document. Even the ones you half-answered. Number them to match the order he asked them.
- **Send the technical note if he wanted one.** Two pages: the pipeline, the four novel decisions, what the code cannot do, and the test numbers from this week.
- **Log the meeting.** Date, what he asked, what he reacted to, what he didn't, what he said the next step was. You will need this in three weeks and you will not remember.

Within two weeks

If he's engaged	If he's quiet
Send a real update — a shipped feature, a user quote, a number that moved. One email, no follow-up question.	One polite nudge at day 7, one at day 14, then stop. "Circling back in case this fell through the cracks — happy to be told it's a no."

If he's engaged	If he's quiet
Ask what would make the partner conversation easy.	Keep the relationship. A no at pre-seed is often a yes at seed if you come back with numbers.
Assume four to eight weeks to a term sheet, and keep raising in parallel until the wire lands.	Never let one firm's silence become your fundraising strategy.

The email itself, drafted

Copy this, cut it to your own words, and have it ready in drafts before the call so that sending it takes one minute rather than forty.

> Subject: Evo — thanks + the two links
>
> Thomas — thanks for the time. The question about [X] was the useful one and I've been thinking about it since.
>
> Three lines as I understood them: Evo reconstructs a workspace from on-device activity evidence; the hard part is deciding which windows are one job, and the rule we use is that production defines identity; the open question for you is whether that becomes a layer or stays a utility.
>
> Two links: the ninety-second recording [link], and the technical note [link]. I've answered your storage question at the bottom of the note — you were right that I didn't have the number.
>
> Next step from your side was [what he actually said]. Happy to work to your timeline.
>
> Bhavit

— The template — one screen, no attachment except one

WHY THIS WORKS

Three things make this email work. It quotes his own question back, so he knows you were listening. It states the open question in his framing, which is what he will actually paste into the memo. And it answers the one thing he caught you on, which is the single highest-signal move available in a follow-up.

19 If it's a no, and if it's a yes

If it's a no

It probably will be, and that is the base rate — not because Evo is bad, but because most first meetings with most firms are no. What matters is the shape of the no, and there are three.

The no	What it means	What you do
"Too early for us."	Usually means "come back at seed." This is the best no available.	Ask what specific milestone would make it a yes. Then go hit it. Email them the day you hit it.
"Too small for the fund."	Honest, structural, not about you.	Ask for an intro to a pre-seed fund they respect. A warm intro from an associate is a real asset and costs him nothing.
"We don't believe the category."	The only no worth arguing with — once, and only in writing.	Ask which premise they doubt. Answer it in one page. Do not argue twice.

WHY THIS WORKS

Whatever the no is, ask one question before you hang up: "If you were me, what would you do next?" Associates give real answers to that, and it is often worth more than the meeting.

If it's a yes — or a maybe

A "yes, let's go deeper" is not a yes. Between the first call and money there is diligence, and diligence on a pre-seed company with an AI-built codebase is specific. Have these ready before he asks.

What he'll want	What you need to have
Incorporation documents	Certificate, bylaws, EIN, cap table, founder stock agreements, 83(b) filings with dates.
IP assignment	Every founder and contractor has assigned IP to the company in writing. If you used AI agents, that's fine — but a human must have assigned the output.
The cap table	Who owns what, on a fully diluted basis, including any advisor grants or promises you've made verbally. Verbal equity promises are the most common pre-seed disaster.
A security overview	One page: what's captured, what's stored, where the keys are, what leaves the device (nothing), what the permissions are.
The technical note	Two pages. Pipeline, four novel decisions, known limits, test numbers.
Two references	People who have seen the product. Not family.

CAREFUL

The single most common thing that kills a pre-seed at diligence is a cap table that doesn't match what the founder said on the call. Know your percentages to the decimal before Friday, including anything you've promised verbally to an advisor or a friend.

PRINT THIS. READ IT AT 10:55PM.

20 The one-page card

The four sentences

1. "Every knowledge worker has a second job they don't get paid for: reconstructing what they were doing before they were interrupted — and it isn't a search problem."

2. "Evo watches work happen locally, groups it into bodies of work, and rebuilds the workspace when you say continue. Sixteen crates of Rust, no HTTP client in it."

3. "The engine has to be honest — every restore reports what it opened, what it couldn't, and there's a corpus of a real working day that fails the build if two works ever merge."

4. "Every agent is getting memory. Almost none are getting jurisdiction — knowing which of your things belong to which job."

The four honest concessions

Traction: verbal, no list, 19-of-26 was an earlier build. **Code:** directed, not hand-written; first hire owns it. **Learning:** compounds memory, not parameters. **Size:** the Mac utility alone is a $5–40M business; the layer is the reason to care.

The four questions to ask him

What happens next and on what timeline. What would make this a yes. What's your biggest concern. How does NEA decide on a round this size.

The physical checklist

Before	Check
T-minus 60 min	Eat. Water. Laptop charged. Phone on silent, face down, out of reach.
T-minus 30 min	Daemon running, three works visible, production created in each, notifications off, desktop clean. Recording tab open.
T-minus 10 min	Zoom audio and screen share tested with a recording, not just opened. Backup hotspot connected and ready. Second device on the same call as a fallback mic.
T-minus 2 min	One sheet of paper, one pen. Section 20 open. Timer visible but not shared.
During	Every unanswerable question goes on the paper. Note the clock at 0:25.
T-plus 2 hours	The email is sent. Not drafted. Sent.

CAREFUL

You are taking this call at 11:30pm IST. The single most common way a call like this goes wrong is not a bad answer — it is a tired founder at midnight, on a bad connection, with a laptop at 8% battery. Treat the logistics as seriously as the pitch.

DRILL

21 Numbers you must know cold

Number	What it is	Say it how
$1.5M	The raise	On a $12M post-money SAFE.
18 months	Runway	At $83K a month, ~$70K of it base.
12.5%	Dilution	1.5 ÷ 12. Don't volunteer it unless asked.
16 crates	Rust workspace	About 138,500 lines.
~1,270	Tests	Re-run before Friday. 59 suites, 0 failed.
10,392	Events in the corpus	One real day. Two days of activity, actually.
164	Hand-labelled resources	Across six bodies of work.
6.1M	US computer & mathematical professionals	BLS 2024. The wedge.
100K	Users at month 12	1.6% of the wedge.
~$1M	ARR run-rate at month 12	9% paying × $108.
$108/yr	Pro price	"The hypothesis."
25 min 26 s	Return-to-task time	Mark et al., CHI 2005. Not a reconstruction stopwatch.
~1,200	App toggles per day	HBR 2022.
Aug 13 2026	OpenAI Computer History launch	The category validation and the threat, in one date.
$311,071	IER investment threshold	The round clears it ~5×.

If he asks about numbers you don't have

There are five numbers an associate may ask for that you genuinely cannot produce. Have the shape of each answer ready, because an improvised number is worse than a missing one.

He asks	You don't have it because	What you say
Cost to acquire a user	You've spent nothing on acquisition.	"Zero spent, so no CAC. The launch budget is precisely to find out, and I'd rather report the first real number than model one."
Activation rate	No instrumentation.	"Not measured yet. Instrumenting the next thousand users from day one is the first engineering task after the round."
Paid conversion	Pro doesn't exist until month three.	"Undefined — there's no paid tier to convert into. I know that's an unsatisfying answer and I'm not going to dress it up."
Storage per user	Not benchmarked across a long history.	"Order of megabytes a day, and I'd rather measure it over a week on my own machine than guess at a decimal."
Support cost	No users at scale.	"Unknown, and I'd expect it to be the real cost driver at a $9 price point rather than infrastructure. That's a hypothesis, not a number."

THE LIST

22 Do not say

Fourteen phrases that cost founders meetings. Some are dishonest, some are just signals of inexperience. All of them are avoidable.

Never say	Because	Say instead
"We have 300 users."	You don't. It's conversations.	"Three hundred conversations, no list."
"19 of 26 retained."	Different build, informal measurement.	"An earlier build, measured by asking, not instrumenting."
"We're the ChatGPT of work."	Category-by-analogy is the mark of a founder who hasn't decided what he's building.	The one-sentence description.
"Nobody else is doing this."	Verifiably false and it destroys credibility in one sentence.	"Everyone remembers. Almost nobody restores."
"It learns and gets better automatically."	There's no fitted model. He will ask what the update rule is.	"It compounds memory, not parameters."
"We have CI on every commit."	There's no .github directory.	"It's a hard test in the workspace suite; the CLI gate exits non-zero."
"The market is $1 trillion."	It's not a market, it's a population.	The bottom-up: 6.1M, 1.6%, 9%, $970K.
"It's end-to-end encrypted."	It isn't yet. Per-record encryption is roadmap.	"FileVault and the sandbox today; per-record keys before sync."
"We're moving to San Francisco."	Puts a government agency on your critical path.	"The US is our market before it's our address."
"Our AI is proprietary."	Meaningless. Everything is proprietary or nothing is.	"The rules are ours; the embedding model is off the shelf."
"We just need a technical co-founder."	Signals you know what you lack and have no plan.	"The first hire is a senior Rust/macOS engineer; here's what they own."
"If we build it, they will come."	No plan.	The three channels and the assumption behind each.
"Our competitor is bad."	Attacking competitors reads insecure.	"They answer a different question."
"Can you sign an NDA?"	Venture firms don't, and it ends conversations.	Nothing. Never ask.

23 Sources and confidence

Everything factual here is graded. **A** means a primary source or the repo. **B** means a reputable secondary source or a reasonable inference. **C** means directional — useful, but don't quote it as fact in the room.

About Thomas and NEA

Claim	Confidenc	Source
Associate, Technology, joined 2025, seed/early and growth	A	nea.com/team/thomas-joshi
Columbia B.Sc. AI; co-author of DSPy; ex-AI engineer in healthcare conversational AI	A	Same profile
Portfolio: Anyscale 2019, Metabase 2019, Factory 2025; Namespace Labs Series A; CuspAI follow-on	A	Profile, NEA blog, X posts
The five quoted lines about evaluation, utility and craft	A	Profile, verbatim or near-verbatim
NEA ~$38B AUM as of 30 June 2026	A	NEA's own fellowship site
NEA 19 ~$5.5B target, first close around Dec 2025–Jan 2026	B	Dealroom summary; not confirmed by NEA directly
NEA seed rounds average ~$4.4M across 241 deals	B	Tracxn
Monday partnership meetings and YE/YS/NO voting	C	Fortune, 2011 — old, and practice will have evolved. Treat as cultural texture, not procedure.
NEA fellowship cohort working on an AI memory thesis	A	NEA's own 2026 fellows blog post

About the market and competition

Claim	Confiden	Source
OpenAI Computer History launched 13 August 2026, macOS, opt-in, Pro/Business/Enterprise, interaction events not screenshots, processed server-side, memories stored locally as Markdown	A	OpenAI documentation and changelog; corroborated by press
Computer History extended to EEA, Switzerland and UK in Aug–Sept 2026	A	OpenAI changelog
Microsoft Recall: Copilot+ only, opt-in, local, encrypted	A	Microsoft documentation
Rewind → Limitless → acquired by Meta Dec 2025; Mac app discontinued	B	Press reporting
NEA led Rewind's $12M at a $350M valuation in 2023	B	Contemporaneous press. Do not raise this in the meeting.
Median post-money SAFE cap $18M for $1M–$2.4M rounds (Q2 2026)	B	Carta data via secondary summaries
~93% of US pre-seed rounds use SAFEs	B	Carta via secondary
Mark et al. 2005: 25 min 26 s average return-to-task, 24 workers	A	CHI 2005 paper

About the code

All **A** — read directly from the repository at commit `7af333bb` .

- 16 crates, ~138,507 lines of Rust across `crates/` /src/ **/** .rs `.`
- `docs/BUILD-LOG.md` : 57 suites / 1,087 tests (30 Sep); `docs/architecture/LOCAL-WORK-RESTORATION.md` : 59 suites / 1,272 tests (5 Oct).
- `crates/evo-threads/tests/ground_truth.rs` : the three hard assertions.
- `tools/threads-regression/` : 10,392 events, 164 labels.
- `crates/evo-execution/src/restore.rs` : opened / raised / unopened-with-reason.
- No `reqwest` , `hyper` , `ureq` or websocket dependency anywhere in the workspace.
- No `.github/` directory — there is no CI configuration.
- `docs/RELEASE-CHECKLIST.md` (6 Oct): end-to-end capture→daemon→UI, cold Accessibility probe, and a real browser profile receipt remain **unpassed** .

> **CAREFUL**
> The last one is the most important line in this book. If he asks whether the product is release-ready, the honest answer is: "It runs and I use it. What I don't have is a signed, notarised build with an end-to-end receipt on a second machine — and the repo documents exactly which gates are unpassed." That answer is far stronger than a confident yes, because it proves you know the difference between a demo and a release.

One last thing. You got this meeting by writing one specific email to one specific person about something she had written. That is the same skill this call requires: say the true, specific thing to the person in front of you, and stop talking. You already know how to do it.