PRIVATE FOUNDER PREPARATION / 06 NOVEMBER 2026

Evo × WMF

The meeting study guide

An investor-style pitch, product walkthrough, agent architecture, hard questions, and a practical plan for a 30–40 minute conversation with Tommaso Suzzi Barberini.

FOR BHAVIT & MANAVV

The outcome to aim for: Tommaso understands the company, has seen an honest demonstration, and can explain exactly which program Evo has been selected for and what happens next.

Do not send this private guide as a pitch deck. Rehearse it. Speak naturally; do not read it aloud.

Research and code inspection: 5 Nov 2026. No macOS runtime verification was performed for this guide. The meeting time and exact WMF call were not supplied.

01 / READ THIS FIRST

Your first 15 minutes of preparation

The mistake to avoid: in the NEA meeting, you did not make the pitch. Tomorrow, do not wait for Tommaso to discover the business through questions. Ask for permission, deliver the pitch, then invite the conversation.

> *"Thank you for selecting Evo. Before we dive in, would it be helpful if we give you a two-minute overview and then show you what we have built?"*

Read in this order

- **First:** memorize the one-line explanation and 2–3 minute pitch on pages 4–5.
- **Next:** rehearse the demo and fallback on page 6; practice the agent example on pages 9–10.
- **Then:** review the red-line claim ledger, the hardest objections, and the WMF questions.

Six anchors

- Today: a macOS MVP helps a person get back to a piece of work. An early query prototype can answer questions about information made available to it.
- Future: user-authorized context for agents and potentially AI vendors. Full agent mode and general vendor integration are not yet built.
- Privacy constraint: no screenshots or screen recordings, even temporarily. That is a design rule—not a full security certification.
- The vision is global: work, study, personal projects and AI users. Start with a narrow repeatable use case; do not confuse the first users with the entire market.
- 300+ conversations are not a waitlist. The 26-person hackathon and 22/26 and 19/26 follow-ups were self-reported, verbal and uninstrumented.
- A WMF selection is not an investment, a competition win, a guarantee of funding, or necessarily a free program.

Who Tommaso is and what this meeting is

Known from your invitation: WMF says Evo has been selected for its Startup Program. Tommaso wants to learn more about the project and walk you through next steps. The meeting is 6 November 2026, 30–40 minutes; the time and format were not provided.

Publicly verified: WMF lists Tommaso Suzzi Barberini with Search On Media Group's innovation team and identifies him with the 2026 Startup District program kickoff. Do not assume he is an investor, the final decision-maker, or personally judging a specific competition. [1]

Unknown: the exact call, program edition, what selection includes, whether it is final, event location, participation costs, services, or obligations. The public Global Startup Program describes visibility, the Startup District and networking, but the invitation does not establish that Evo is in that exact program. [2]

Some WMF competition pages describe a 3-minute pitch and 2-minute Q&A.; That is useful rehearsal practice, **not confirmation that tomorrow's conversation has this format.** The AI for Future page contains conflicting dates; ask Tommaso rather than guessing. [3]

Suggested 35-minute structure — adapt to his agenda

- 0–3 min: greetings, confirm his agenda, ask permission for a short pitch.
- 3–6 min: investor-style pitch: problem → product today → agent vision → evidence → business → ask.
- 6–10 min: show a small, safe, rehearsed demo. If he wants questions first, follow his lead.
- 10–25 min: answer product, team, business, privacy and competition questions.
- 25–32 min: ask about exact program, value, costs, selection and next steps.
- 32–35 min: agree owner, documents and date for follow-up. If 40 minutes, leave room for questions.

Handoff: Bhavit leads problem, product, technical answers and demo. Manavv leads market, business, program questions and next-step summary. Either founder can answer naturally; do not force a scripted handoff.

The one sentence and 30-second version

One sentence

> *“Evo helps people get back to the work they were doing, and we are building toward making that context useful to the AI agents they choose.”*

Thirty seconds

> *“Our work is spread across documents, websites and apps. When we return to a project, we have to find the pieces again; when we ask an AI agent for help, we often have to explain the context again. Evo starts by helping someone return to a piece of work on their Mac. We are also exploring how an agent could use relevant context, with the person’s permission. We have an early MVP and query prototype; full agent mode is still being built.”*

The bridge from continuity to agents

The two ideas are related but they are not the same product maturity. Returning to work is the human-facing use case. An agent-facing interface is a proposed next step. A context record may contain resource names and links; it cannot magically know unseen decisions or unfinished intentions.

Language to keep straight

- **Work:** any activity someone wants to continue, including study, a trip, a personal project or employment.
- **Context:** the relevant, available information about that activity—not the person’s complete memory or all content on their computer.
- **Agent:** software using a model and tools to help carry out a task. Evo is not the agent; the longer-term idea is that Evo can supply context to one.
- **One action:** the present value is returning to work. A future agent can query context; acting in other applications is a separate, higher-risk capability.

04 / MAIN PITCH

Your two- to three-minute investor-style pitch

Bhavit: “Thank you for inviting us. We are building Evo because people’s activities are scattered across their computers. A project might involve a document, a few websites and an application. When you come back after an interruption, the information is still there, but you have to find the pieces and remember how they relate.”

“Now there is a second version of the same problem. If you ask an AI agent to help with that project, you often have to explain the context again. The model may be capable, but it does not automatically know which activity you mean or which resources are relevant.”

“Evo starts with the person. Our current macOS MVP can help someone get back into a piece of work. Underneath, the system uses available activity signals to form a view of work and the associated resources. It does not take screenshots or screen recordings. We have also built an early prototype that can answer questions about information made available to it.”

“The longer-term vision is a personal context system: a person can continue their own activity, and eventually choose to make relevant context available to an AI agent. We are not claiming that full agent mode or a general vendor integration is already built.”

Manavv: “The long-term market is global: people using computers for professional work, education and personal projects, and eventually the AI products they use. Our practical starting point is people who repeatedly move between projects and apps on a Mac. We need to prove they return, get value and are willing to pay before expanding.”

“We see two potential business models: an individual subscription and, later, a commercial context service for AI developers or vendors. Both pricing and vendor willingness to pay remain to be validated.”

“Our earlier MVP was tested by 26 people at a college hackathon. Follow-up was verbal and self-reported, so we do not treat that as instrumented retention. We are now focused on a more reliable product, repeat-use measurement and testing an agent integration.”

“We are looking for relevant investor and product connections, and we would love to understand the exact next steps and opportunities in the WMF Startup Program. Bhavit can show you what works today.”

If asked about fundraising: “We are seeking pre-seed funding. We are evaluating the right amount against a detailed product and operating plan.” The earlier $1M was a proposed planning assumption, not a confirmed ask. Do not state a valuation or committed runway.

A calm, credible demonstration

Before the call

- Rehearse on the actual Mac and exact build. Verify permissions, startup, the resource you intend to open and the behavior you plan to show. Do not assume repository code works end to end.
- Use a synthetic, non-sensitive activity: one test document and a safe browser URL. Close personal messages, bookmarks, confidential tabs, filenames and notifications before screen sharing.
- Confirm the internet connection, screen-sharing permission and the backup plan. Screen sharing for this meeting is distinct from Evo capturing the screen: Evo itself must not capture screenshots or recordings.

The 90-second demo script

> *"I'll show the part that exists today. This is a test activity with a document and a web resource. I switch away from it. Now I ask Evo to bring me back to that piece of work. Here is what the current build identifies, and here is what it can actually reopen."*

Narrate only the visible output. Distinguish opening a URL from restoring the original tab; opening a document from recreating the exact editor state. If an item fails, say so. Avoid implying the entire work environment was reconstructed.

If showing the query prototype

> *"This is a separate early query prototype, not our completed agent integration. It answers using the information made available to it. The next step is making a narrow, user-authorized version reliable for an external agent."*

If the demonstration uses the local boundary service, know whether its input is real engine output or a label-derived fallback. The service can use a fallback file if engine output is missing; do not call fallback data a live inference. Do not show the real private event corpus. [4]

If the demo fails

> *"The live build is not behaving as expected, so I do not want to imply it is. Let me explain the workflow and we can send a verified demonstration afterwards."*

Do not use a fake recording, invented numbers, or private data to recover the moment.

From activity to a return suggestion

Plain-English sequence: Evo observes certain supported, structured events; stores or processes those events; derives candidate activities and associated resources; and presents a way to return to supported resources. Which collectors, permissions, views and restoration paths work in the current Mac build must be checked before the demo.

1. Observe — not screen images

Repository code defines structured macOS signals such as window focus, file saves, URL navigation and application activation. Some paths also define content-free aggregated input counts. An event type in code is not proof that it is active or wired into your current build. Names, paths, URLs and titles can themselves be sensitive. [5]

2. Derive a candidate activity

The Rust engine replays observed evidence to derive work groupings and a return-oriented view. It uses evidence and deterministic rules rather than a claim that an AI model perfectly reads intent. Some resources can be ambiguous or unassigned. Shared resources can belong to multiple activities; the difficult problem is avoiding false merges.

3. Present and return

A view can associate a candidate activity with resources. Reopening a known resource is an operation on that resource, not proof of complete state restoration. The exact supported resources and action path are build-dependent. Test them on your Mac.

What metadata cannot tell you

A document title or URL may tell Evo what was open; it does not reveal the user's private reasoning, a decision made in a meeting, a specific unfinished sentence or a task deadline. A better answer may require an explicit user note or a separately authorized source. Do not invent that content.

> *"Evo works from evidence it can actually access. When the evidence is insufficient, the honest result is uncertainty, not a confident story about what the person intended."*

What is implemented, prototyped and proposed

A. Source-level observations — not a runtime certification

- The macOS adapter defines focus, navigation, file-save and other structured event types. The thread engine includes work grouping, derivation and a resume-oriented representation. The current build and actual capture path have not been tested here. [5][6]
- The repository includes a local Python prototype with a query interface and an MCP-over-stdio mode. It reads either exported engine output or a label-derived fallback. That demonstrates a possible integration shape; it does not prove a production agent service. [4]
- The prototype contains deterministic query matching and scenario-specific keyword expansion. Do not describe it as a general reasoning model or claim it understands arbitrary tasks.

B. Specific caveats

- The engine has concepts for explicit declarations, but the reviewed prototype does not expose declaration writes; example gate buttons are simulated. Do not demonstrate a simulated change as a persisted permission or learning event. [7]
- Prior materials mention an example where the engine claimed 2 of 8 hand-labelled works. That is one internal example, not an accuracy or coverage benchmark. Do not generalize it or show private material.
- A comment, a proposed design, a saved fixture, or a historical test report is not proof of a passing live demo, security guarantee or a generalized agent product.

C. Unbuilt / unproven

Production-ready agent integration; authenticated per-user/work-scoped authorization; comprehensive privacy and security controls; arbitrary app access; reliable Windows/mobile/cross-device sync; general agent actions; mature commercial API; real-world accuracy at scale.

> *“We have a useful early human-facing MVP and a separate query prototype. The work ahead is turning evidence into reliable context, exposing it safely, and validating whether that improves real tasks.”*

How an agent would actually use Evo

The simplest answer: “The agent asks Evo for relevant context about the activity the person has chosen. Evo returns a small, supported set of context, and the agent uses it to help. That interface and its permissions are what we need to build and validate.”

Example: prepare a project update — future flow, not a shipped demo

- **1. Person:** “Help me prepare an update for my thesis project.” They select the activity and choose which AI product to connect.
- **2. Agent:** asks an Evo tool for context related to that specific activity. The request is not a request for the whole computer.
- **3. Evo:** checks the requesting client and the person’s authorization, finds candidate resources, and returns only permitted, relevant evidence. If confidence is low, it asks for clarification or returns insufficient evidence.
- **4. Agent:** uses the returned titles, resource references and any separately authorized content to formulate an answer. It should distinguish an actual note from an inference.
- **5. Person:** reviews the output. If the agent needs to read source content or take an action, the applicable additional access and approval are separate questions.

A realistic boundary

If Evo only knows that a thesis file and research URLs were open, it can help identify those resources. It cannot say “you decided to use method X” unless an authorized source actually supplies that fact. An agent must not turn a filename into a fabricated conclusion.

Local versus remote

A local agent can potentially speak to a local Evo process. A remote AI service needs a separately engineered, authenticated bridge and an explicit decision about what leaves the device. “Local-first” does not mean data remains local after a user chooses to send it to a remote model. No such production bridge is claimed today.

What the interface could look like

This is a **proposed design**, not an existing production API. It is here so you can explain the mechanism, boundaries and decisions—not claim that implementation is finished.

Narrow read-only tools to explore

- **find_activity(query):** returns candidate activity labels and opaque IDs, not the entire history. Ask the person to choose if results are ambiguous.
- **get_activity_context(activity_id):** returns a small list of authorized resource descriptors: title or label, type, last observed time and provenance. If authorized, a short explicit user note could be included; do not assume notes exist.
- **get_resource(reference):** only if separately permitted and actually supported; returning a reference does not imply the agent can read the file's contents.

An illustrative response

Activity: "Thesis project"
Evidence: document "Thesis outline" (last observed), research URL (last observed).
Unknown: the person's latest decision and next task.
Permission: illustrative only; production authorization is not yet implemented.

MCP and API explained

MCP is a protocol through which an agent host discovers and calls tools. Evo could expose a limited local MCP tool or a future authenticated remote API. MCP is the language of the request; it does not itself guarantee data safety, correctness or permission enforcement. OpenAI documents external MCP tools and warns about prompt injection. [8][9]

Actions require another boundary

A context lookup is read-only. Opening files, editing, sending messages or running commands requires separate policies and user approval. An answer that says "only this work" does not sandbox a process. OS-level containment and adversarial tests would be required for actual enforcement.

How to answer privacy questions

Do you take screenshots?

“No. Our product direction rules out screenshot and screen-recording capture, even transiently. We instead work with structured signals. Those signals can still be sensitive, so we cannot claim that the design choice alone solves privacy.”

What data does Evo collect?

“The code defines supported signals such as app or window focus, file and URL events and timestamps. What is enabled and captured in the actual build must be verified. We want to minimize collection and clearly disclose each source and permission.” Do not list optional inputs as definitely active.

Can a vendor access everything?

“That is not the intended design. We want a person to authorize a particular connection and limit what it can receive. This is a proposed control model, not a claim that broad production-grade enforcement already exists.”

What if an agent or retrieved content is malicious?

“We should treat external content as untrusted, keep the tool surface narrow, authenticate clients, isolate data, log access appropriately, require approval for consequential actions, and test prompt-injection and exfiltration paths. These are engineering requirements, not a guarantee that we have completed them.” [8][9]

Can a user revoke access?

“Revocable access is part of the design we intend to implement and test. We cannot claim that production revocation is complete today.” Revocation can stop future access; it cannot guarantee a third party deletes information already received.

Is Evo GDPR-compliant?

“We have not completed the legal and technical assessment needed to make that claim. For European deployment we need to address lawful basis, transparency, minimization, retention, deletion, international transfers and security with qualified counsel.” Consent is not automatically the correct legal basis, especially in employment contexts. [10]

Privacy by design is a direction; compliance and security are evidence-backed work, not a slogan.

What you can say about progress

Traction — use the precise labels

> *“We have had more than 300 conversations about Evo. That is not a waitlist or a count of active users. We have also had over 50 interviews, including potential-user and cofounder or networking conversations; there may be overlap.”*

> *“An earlier MVP was tested by 26 people at a college hackathon. On verbal follow-up, the founders report that 22 said they were still using it at day 7 and 19 at day 30. That is a small, self-reported and uninstrumented cohort—not verified product analytics.”*

Do not imply a larger sample, organic demand, conversion, independent verification or product-market fit. If the exact interview questions and follow-up records are not available, lead with the 26 testers and omit the percentages. Earlier deck waitlist and testimonials should not be repeated as facts.

Commercial model

Initial hypothesis: give individuals a useful continuity product, then test a paid tier after repeated value is clear. Later, explore developer/vendor pricing for context delivery—possibly usage-based, per authorized connection or negotiated contracts. No vendor revenue or willingness to pay is established.

What makes this venture-scale?

> *“The first product solves a recurring problem for a person. If many people find it valuable and choose to connect it to their AI tools, a vendor-facing context service may become a second business. The global opportunity is large, but distribution, trust, willingness to pay and platform competition are what we must prove.”*

Pricing answer

> *“We have not validated pricing. We are considering an individual subscription and a later vendor service, but we want to learn from actual use and buyer conversations before committing to a price.”*

Global vision, clear first market

Who is Evo for?

“Ultimately, people across the world using their digital systems—students, professionals and people working on personal projects. Agents and AI vendors are also potential beneficiaries or customers. We start with Mac users who repeatedly return to activities across apps, because we can test that use case now.”

How big is the market?

“The broader digital and AI populations are very large, but not all of them are reachable or paying customers. We have built illustrative individual and vendor models; the assumptions are not validated. We are focusing first on adoption and buyer willingness to pay.” If he wants details, see page 15.

Is the market only developers?

“No. Developers may be an accessible early group, but the product is about activities, not occupation. A student returning to research and a professional returning to a project can have the same continuity problem.”

Who is the initial buyer?

“Our initial hypothesis is an individual who experiences the problem regularly and can choose to install the product. Whether that person pays, or an employer or institution eventually pays, still needs testing.”

How do you acquire users?

“Founder-led outreach and communities around Mac productivity, builders, students and recurring projects; then focus on real activation and repeat use. We are not claiming an established channel or partnership.”

Why would AI vendors buy?

“Potentially to improve an agent task using relevant context the user authorizes. We must show measurable task improvement, sufficient authorized users, a reasonable integration cost and economics better than building or buying elsewhere. We have not validated their willingness to pay.”

Is this a data broker?

“Our intended model is a user-directed context service, not selling unrestricted access to a person’s history. The exact data flows and commercial terms still need to be designed and legally reviewed.”

Answer directly; do not pretend nobody else exists

Isn’t Recall already doing this?

“Recall is a direct competitor in return-to-work. Its public documentation says it periodically reads screen images; our design excludes screenshot and screen-recording capture entirely. Our long-term interest is also user-directed context for agents. We have to prove our approach is useful; the no-image choice alone is not a moat.” [11]

Can ChatGPT or Apple build this?

“Yes. Major platforms have distribution and are already building memory and context features. OpenAI also documents an optional Computer History feature for macOS enterprise users. Our opportunity is to make a person-centered experience useful across activities and, eventually, tools—but cross-vendor value and integrations have to be earned.” [12]

How is Evo different from Mem0 or Supermemory?

“They provide developer-facing memory infrastructure. Evo starts from a person’s activities and the continuity problem; we want to explore making that context available to agents. The distinction is our approach, not a claim that they cannot enter this space or that our integration is already as mature.” [13]

What is defensible?

“Potential defensibility would come from a product people trust and use repeatedly, high-quality context derived from actual usage, integrations that improve agent outcomes and a distribution advantage. Today these are things to build and measure, not proven moats.”

Why not just use bookmarks or folders?

“People can, and for some they may be enough. Evo must reduce the effort of reconstructing an activity enough that users prefer it. We need to compare against those real alternatives, not assume they fail.”

Market model: explain it without treating it as fact

For this screening call, lead with the customer, use case and evidence. Only use these calculations if asked. They are **illustrative long-term scenarios**, not a sourced market report or a revenue forecast.

Individual scenario

6.15B estimated internet users × assumed 20% eventual eligible and paying share × assumed $60/year = $73.8B/year. The user base includes mobile-only users and people Evo may never serve; both the conversion and price are assumptions. It is not today's serviceable market. [14]

Vendor scenario

Menlo estimates 2B consumer AI users globally. If 25% eventually connect, each has 1.5 paid vendor connections, and a vendor pays $0.50 per connection per month, the result is $4.5B/year. Only the 2B is an external estimate; all other inputs are unvalidated. The 2B is not a monthly agent-user base. [15]

Sanity check: the base assumes 500M connected people and 750M billable person–vendor connections. If 25 vendors supplied them all, that averages 30M connections per vendor. This is an ambitious end state, not near-term expected adoption.

Do the revenue streams add?

Only if a person's subscription and a vendor's fee are genuinely separate payments to Evo. If one replaces the other or the vendor passes the fee through to the same customer, adding them can overstate the opportunity. Do not add the number of individual users and AI users as if they were unique populations.

The honest short answer

> *"We see two potential revenue streams and a global long-term market. Our market models are scenarios, not validated demand. The next step is to prove repeat use and speak to vendors about whether the context materially improves their product and what they would pay."*

Short answers that survive follow-ups

What exactly is built right now?

“A macOS MVP for helping a person return to a piece of work, plus a separate early query prototype. The complete agent product and production vendor integrations are still to be built.” Describe only what you can demonstrate on your machine.

Does Evo learn what I am doing?

“It derives candidate work from evidence using rules. We should not say the current engine has learned your intentions or that it trains itself from every correction. We are investigating how user feedback could improve the experience.”

How do you group work accurately?

“We use available structured activity evidence to derive candidate groupings. A shared application is not enough to prove two activities are the same. Ambiguity and mistaken merges are real problems, so we need labeled, multi-user evaluation and a way to surface uncertainty.”

What if it groups a personal task with a client project?

“That would be a serious relevance and privacy failure. We want to minimize exposure, allow corrections and test false merges specifically. A future external interface must not share data based only on an unverified grouping.” Do not claim controls already enforce this.

What if someone asks a question and the information is not there?

“The system should say it lacks enough evidence. Metadata does not give us a person’s unspoken intent. If additional information is needed, the person can provide it or separately authorize a source in a future product.”

Is it an AI product if the current grouping is rule-based?

“The current engine uses deterministic logic for its grouping. The AI opportunity is to help an agent work with the right context. We describe the current technology accurately rather than calling every part AI.”

How do you know it saves time?

“We believe it can reduce the effort of finding resources, but we have not established a verified time-saving figure. We plan to compare real return tasks with the tools people already use.”

Product, business and execution

Can I use it on Windows or my phone?

“The MVP we are discussing is on macOS. We want to reach more devices, but Windows, mobile and sync are not verified shipped capabilities in this build.”

Does it work offline?

“The engine has local components, but we need to verify the exact build and every service used before promising fully offline operation. If a person chooses a remote AI service, the permitted data necessarily travels to that service.”

Are you training on people’s data?

“We should answer based on the actual build and current policies. The reviewed query prototype is not proof of an AI training pipeline; we should not make a blanket company-wide guarantee without verifying every component.”

What are you raising?

“We are seeking pre-seed funding and are building a detailed operating plan for product, security and distribution. An earlier working model was $1M for roughly 18 months, but that is not a finalized budget or guaranteed runway.” Use an approved ask if you have one.

Are you incorporated? Where is the team?

“We are currently in India.” State the legal entity and incorporation status only after the founders verify it. Manavv has a UK post-study work permit; do not promise a return date or claim it gives permission to work elsewhere.

Why this team?

Bhavit: “I built the Rust/SQLite engine and experienced the problem myself.” Manavv: “I focus on the market, positioning and channels.” Mention the prior deck’s Buildspace, Delhi and Leeds credentials only as founder-confirmed details; do not invent exits or revenue.

What do you need next?

“A reliable repeat-use product, a well-designed evaluation, a narrow agent-context prototype, privacy and security groundwork, and the right early-user and investor connections.”

Questions to ask Tommaso

The goal is to understand what you were selected for, whether it is useful, what it costs and what to do next. This is a professional two-way conversation—not an assumption that WMF is asking you for money or offering funding.

Ask these first

- “Could you clarify the exact name and edition of the program we have been selected for, and what this selection means?”
- “Is our selection final for the program, or are there further review stages? Is it different from the Startup Competition or AI for Future call?”
- “What specific activities are included: a pitch, exhibition space, investor matchmaking, mentoring, or something else? Which are guaranteed versus application-based?”
- “What are the participation, booth, ticket, service or administrative fees, if any? Are travel and accommodation separate? Could you share the written terms and cancellation/refund policy?”
- “What are the event dates and location, and is attendance in person mandatory? Are there remote options?”

If the fit is good

- “Which investors, AI companies or potential design partners could be relevant to a personal-context product, and how does matchmaking work?”
- “What would make Evo a strong fit for the ecosystem, and what should we improve before any pitch?”
- “What materials do you need from us, who owns each next step, and by when?”

Public WMF pages list opportunities and evaluation criteria, but benefits differ by call and can involve additional selection. Do not assume finalists, prize amounts, funded travel or introductions from a generic selection email. [2][3]

Finish the call with a concrete next step

A good closing

> *“Thank you, Tommaso. To make sure we understood: we are selected for [his exact program], the next step is [his explanation], and we will send [agreed material] by [date]. Is that right?”*

Then ask for the written program details, fee schedule and point of contact. If something remains unclear, ask before committing. Do not agree to a paid package on the call without reviewing the terms.

Post-call email template

Subject: Thank you — Evo and WMF next steps

Hi Tommaso,

Thank you for taking the time to speak with us. We enjoyed sharing Evo and learning about [exact program].

As discussed, we will send [material] by [date]. Our understanding is that the next step is [step], and you will share [program details / terms / timetable]. Please let us know if we missed anything.

Best,
Bhavit and Manavv

Tomorrow’s pre-call checklist

- Read the invitation and any application copy; identify discrepancies with your current product and prepare to correct them.
- Agree who opens, who demos, who covers the business, and who records follow-up actions.
- Run the actual demonstration twice with synthetic, non-sensitive data; verify what opens and what does not.
- Confirm current product facts, availability, data handling, incorporation and fundraising ask. Do not guess.
- Have a stable connection, the call link, product build, notebook and this guide ready. Leave notifications off.
- Practice the two-minute pitch aloud once each. Stop when the pitch ends and invite questions.

Read this five minutes before the call

Open

> *“Would it be helpful if we give you a two-minute overview and then show you what we have built?”*

Pitch in six beats

1. Activities are scattered across apps and resources.
2. People reconstruct context; agents face a related context problem.
3. Today: macOS return-to-work MVP plus early query prototype.
4. Future: user-authorized context for agents; not yet general integration.
5. Global vision, initial recurring Mac workflows, individual then vendor business hypotheses.
6. Small early test; measure repeat use and payment; seek the right WMF connections.

Three sentences to remember

- “We do not capture screenshots or screen recordings, even temporarily.”
- “The current MVP helps people return; the agent integration is the next thing to build and validate.”
- “We do not know that yet. Here is how we intend to test it.”

Ask Tommaso

Exact program and selection stage? What is included? Any fees? Dates and location? In-person required? Investor and vendor access? Next step, owner and deadline?

Never say

“Only Evo restores”; “our agent works with every vendor”; “we understand everything”; “we never send any data”; “we are GDPR compliant”; “we have 300 waitlist signups”; “73% instrumented retention”; “WMF invested in us.”

Research trail and claim checks

Sources were consulted on 5 November 2026. External pages can change. Program details must be confirmed directly with WMF. Local paths below refer to the private working repository; do not send private corpus files.

[1] WMF: Tommaso and Startup District kickoff
https://en.wemakefuture.it/s/69394a65f3491b86a26cdd3b/

[2] WMF Global Startup Program (possible relevant program; not confirmed for Evo)
https://en.wemakefuture.it/call/global-startup-program/

[3] WMF AI for Future: criteria and pitch format; page contains conflicting dates
https://en.wemakefuture.it/call/startup-ai-for-future/

[4] Local prototype and source/fallback distinction
exports/ship-happens/README.md; exports/ship-happens/evo_server.py

[5] macOS structured event definitions
crates/evo-capture/src/adapters/macos.rs

[6] Static implementation handoff and limitations
docs/research/RECALL-TO-EVO-IMPLEMENTATION-HANDOFF.md

[7] Declaration versus simulated gate
exports/ship-happens/JURISDICTION.md

[8] MCP authorization and security
https://modelcontextprotocol.io/specification/2025-11-25/basic/authorization

[9] OpenAI MCP tools and prompt-injection risks
https://developers.openai.com/api/docs/guides/tools-connectors-mcp

[10] European Data Protection Board: principles and lawful processing
https://www.edpb.europa.eu/sme/learn-the-basics/data-protection-basics_en

[11] Recall public product and privacy statements
https://www.joinrecall.com/

[12] OpenAI enterprise release notes: optional Computer History
https://help.openai.com/en/articles/10128477-chatgpt-enterprise-and-edu-release-notes

[13] AI memory products
https://supermemory.ai/ | https://mem0.ai/

[14] DataReportal global digital overview
https://datareportal.com/global-digital-overview

[15] Menlo 2026 consumer AI methodology
https://menlovc.com/perspective/2026-the-state-of-consumer-ai/